

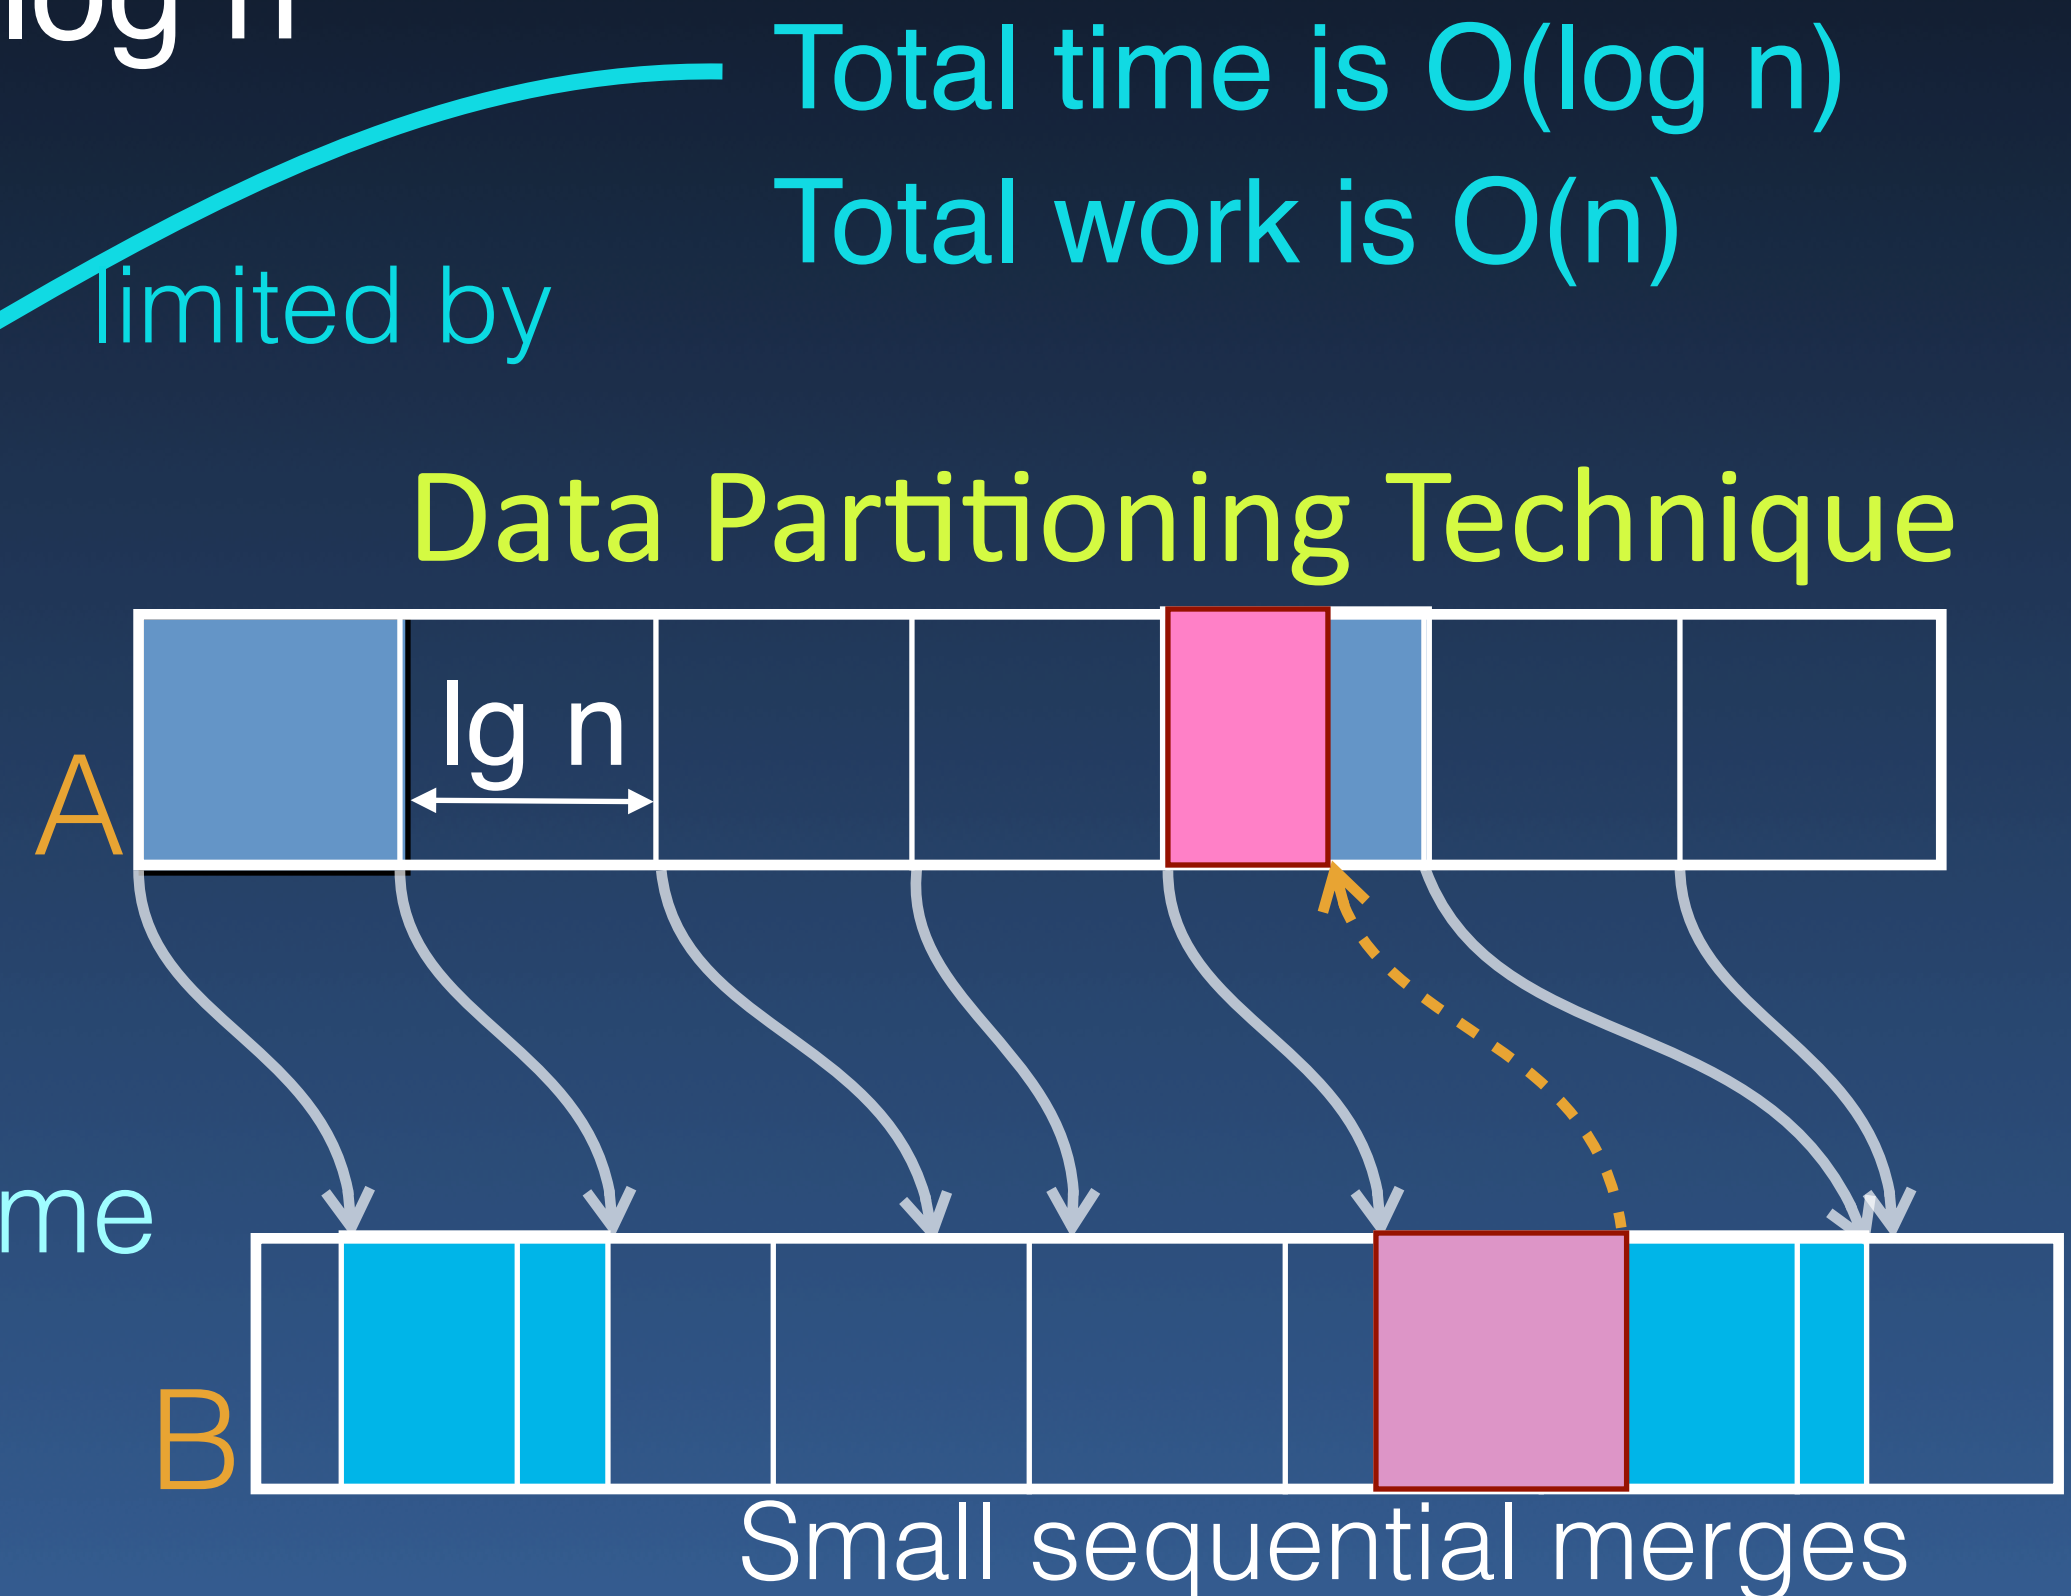
COL380

Introduction to
Parallel & Distributed Programming

- Parallel Algorithms
 - ➔ Merging
 - ➔ Minimum finding
 - ➔ Sorting

Towards Optimal Merge(A,B)

- Partition A and B into $\log n$ sized blocks
- Select from A, elements $i * \log n$, $i \in 0:n/\log n$
- Rank each selected element of A in B
 - Binary search
- Merge pairs of sub-sequences
 - If $|B_i| \leq \log(n)$, Sequential merge in $O(\log n)$ time
 - Otherwise, partition B_i into $\log n$ blocks
 - And Recursively subdivide A_i into sub-sub-sequences



Fast Merge (A,B)


- Partition A and B into \sqrt{n} blocks each
- Select from A, elements $i\sqrt{n}$, $i \in [0: \sqrt{n})$



Fast Merge (A,B)

- Partition A and B into \sqrt{n} blocks each
- Select from A, elements $i\sqrt{n}$, $i \in [0: \sqrt{n})$
- Rank each selected element of A in B

→ \sqrt{n} Parallel searches, use \sqrt{n} processors for each search




Fast Merge (A,B)

- Partition A and B into \sqrt{n} blocks each
- Select from A, elements $i\sqrt{n}$, $i \in [0: \sqrt{n})$
- Rank each selected element of A in B


→ \sqrt{n} Parallel searches, use \sqrt{n} processors for each search

- Similarly rank \sqrt{n} selected elements from B in A



Fast Merge (A,B)

- Partition A and B into \sqrt{n} blocks each
- Select from A, elements $i\sqrt{n}$, $i \in [0: \sqrt{n})$
- Rank each selected element of A in B



→ \sqrt{n} Parallel searches, use \sqrt{n} processors for each search

- Similarly rank \sqrt{n} selected elements from B in A
- Recursively merge pairs of sub-sequences

→ Total time: $T(n) = O(1) + T(\sqrt{n}) = O(\log \log n)$

→ Total work: $W(n) = O(n) + \sqrt{n} W(\sqrt{n}) = O(n \log \log n)$


Fast, but too much work
Not work optimal

Optimal Merge (A,B)

- Use the fast, non-optimal algorithm on small enough subsets
- Subdivide A and B into blocks of size **$\lg n$** ($\lg = \log \log$)

→ A_1, A_2, \dots

→ B_1, B_2, \dots



Optimal Merge (A,B)

- Use the fast, non-optimal algorithm on small enough subsets

- Subdivide A and B into blocks of size $\lg n$ ($\lg = \log \log$)


→ A_1, A_2, \dots

→ B_1, B_2, \dots

- Select first element of each block

→ $A' = p_1, p_2, \dots$

→ $B' = q_1, q_2, \dots$



Optimal Merge (A,B)

- Use the fast, non-optimal algorithm on small enough subsets

- Subdivide A and B into blocks of size $\lg n$ ($\lg = \log \log$)


→ A_1, A_2, \dots

→ B_1, B_2, \dots

- Select first element of each block

→ $A' = p_1, p_2, \dots$

→ $B' = q_1, q_2, \dots$



Optimal Merge (A,B)

- Use the fast, non-optimal algorithm on small enough subsets

- Subdivide A and B into blocks of size $\lg \lg n$ ($\lg = \log \log$)


→ A_1, A_2, \dots

→ B_1, B_2, \dots

- Select first element of each block

→ $A' = p_1, p_2, \dots$

→ $B' = q_1, q_2, \dots$



Optimal Merge (A,B)

- Use the fast, non-optimal algorithm on small enough subsets

- Subdivide A and B into blocks of size $\lg \lg n$ ($\lg = \log \log$)


→ A_1, A_2, \dots

→ B_1, B_2, \dots

- Select first element of each block

→ $A' = p_1, p_2, \dots$

→ $B' = q_1, q_2, \dots$



Optimal Merge (A,B)

- Use the fast, non-optimal algorithm on small enough subsets

- Subdivide A and B into blocks of size $\lg n$ ($\lg = \log \log$)


→ A_1, A_2, \dots

→ B_1, B_2, \dots

- Select first element of each block

→ $A' = p_1, p_2, \dots$

→ $B' = q_1, q_2, \dots$



Optimal Merge (A,B)

- Use the fast, non-optimal algorithm on small enough subsets

- Subdivide A and B into blocks of size $\lg n$ ($\lg = \log \log$)

→ A_1, A_2, \dots


→ B_1, B_2, \dots

- Select first element of each block

→ $A' = p_1, p_2, \dots$

→ $B' = q_1, q_2, \dots$

- Now merge separated blocks of A and B



Optimal Merge (A,B)

- Use the fast, non-optimal algorithm on small enough subsets

- Subdivide A and B into blocks of size $\lg n$ ($\lg = \log \log$)

→ A_1, A_2, \dots


→ B_1, B_2, \dots

- Select first element of each block

→ $A' = p_1, p_2, \dots$

→ $B' = q_1, q_2, \dots$

- Now merge separated blocks of A and B



Optimal Merge (A,B)

- Use the fast, non-optimal algorithm on small enough subsets

- Subdivide A and B into blocks of size $\lg n$ ($\lg = \log \log$)

→ A_1, A_2, \dots

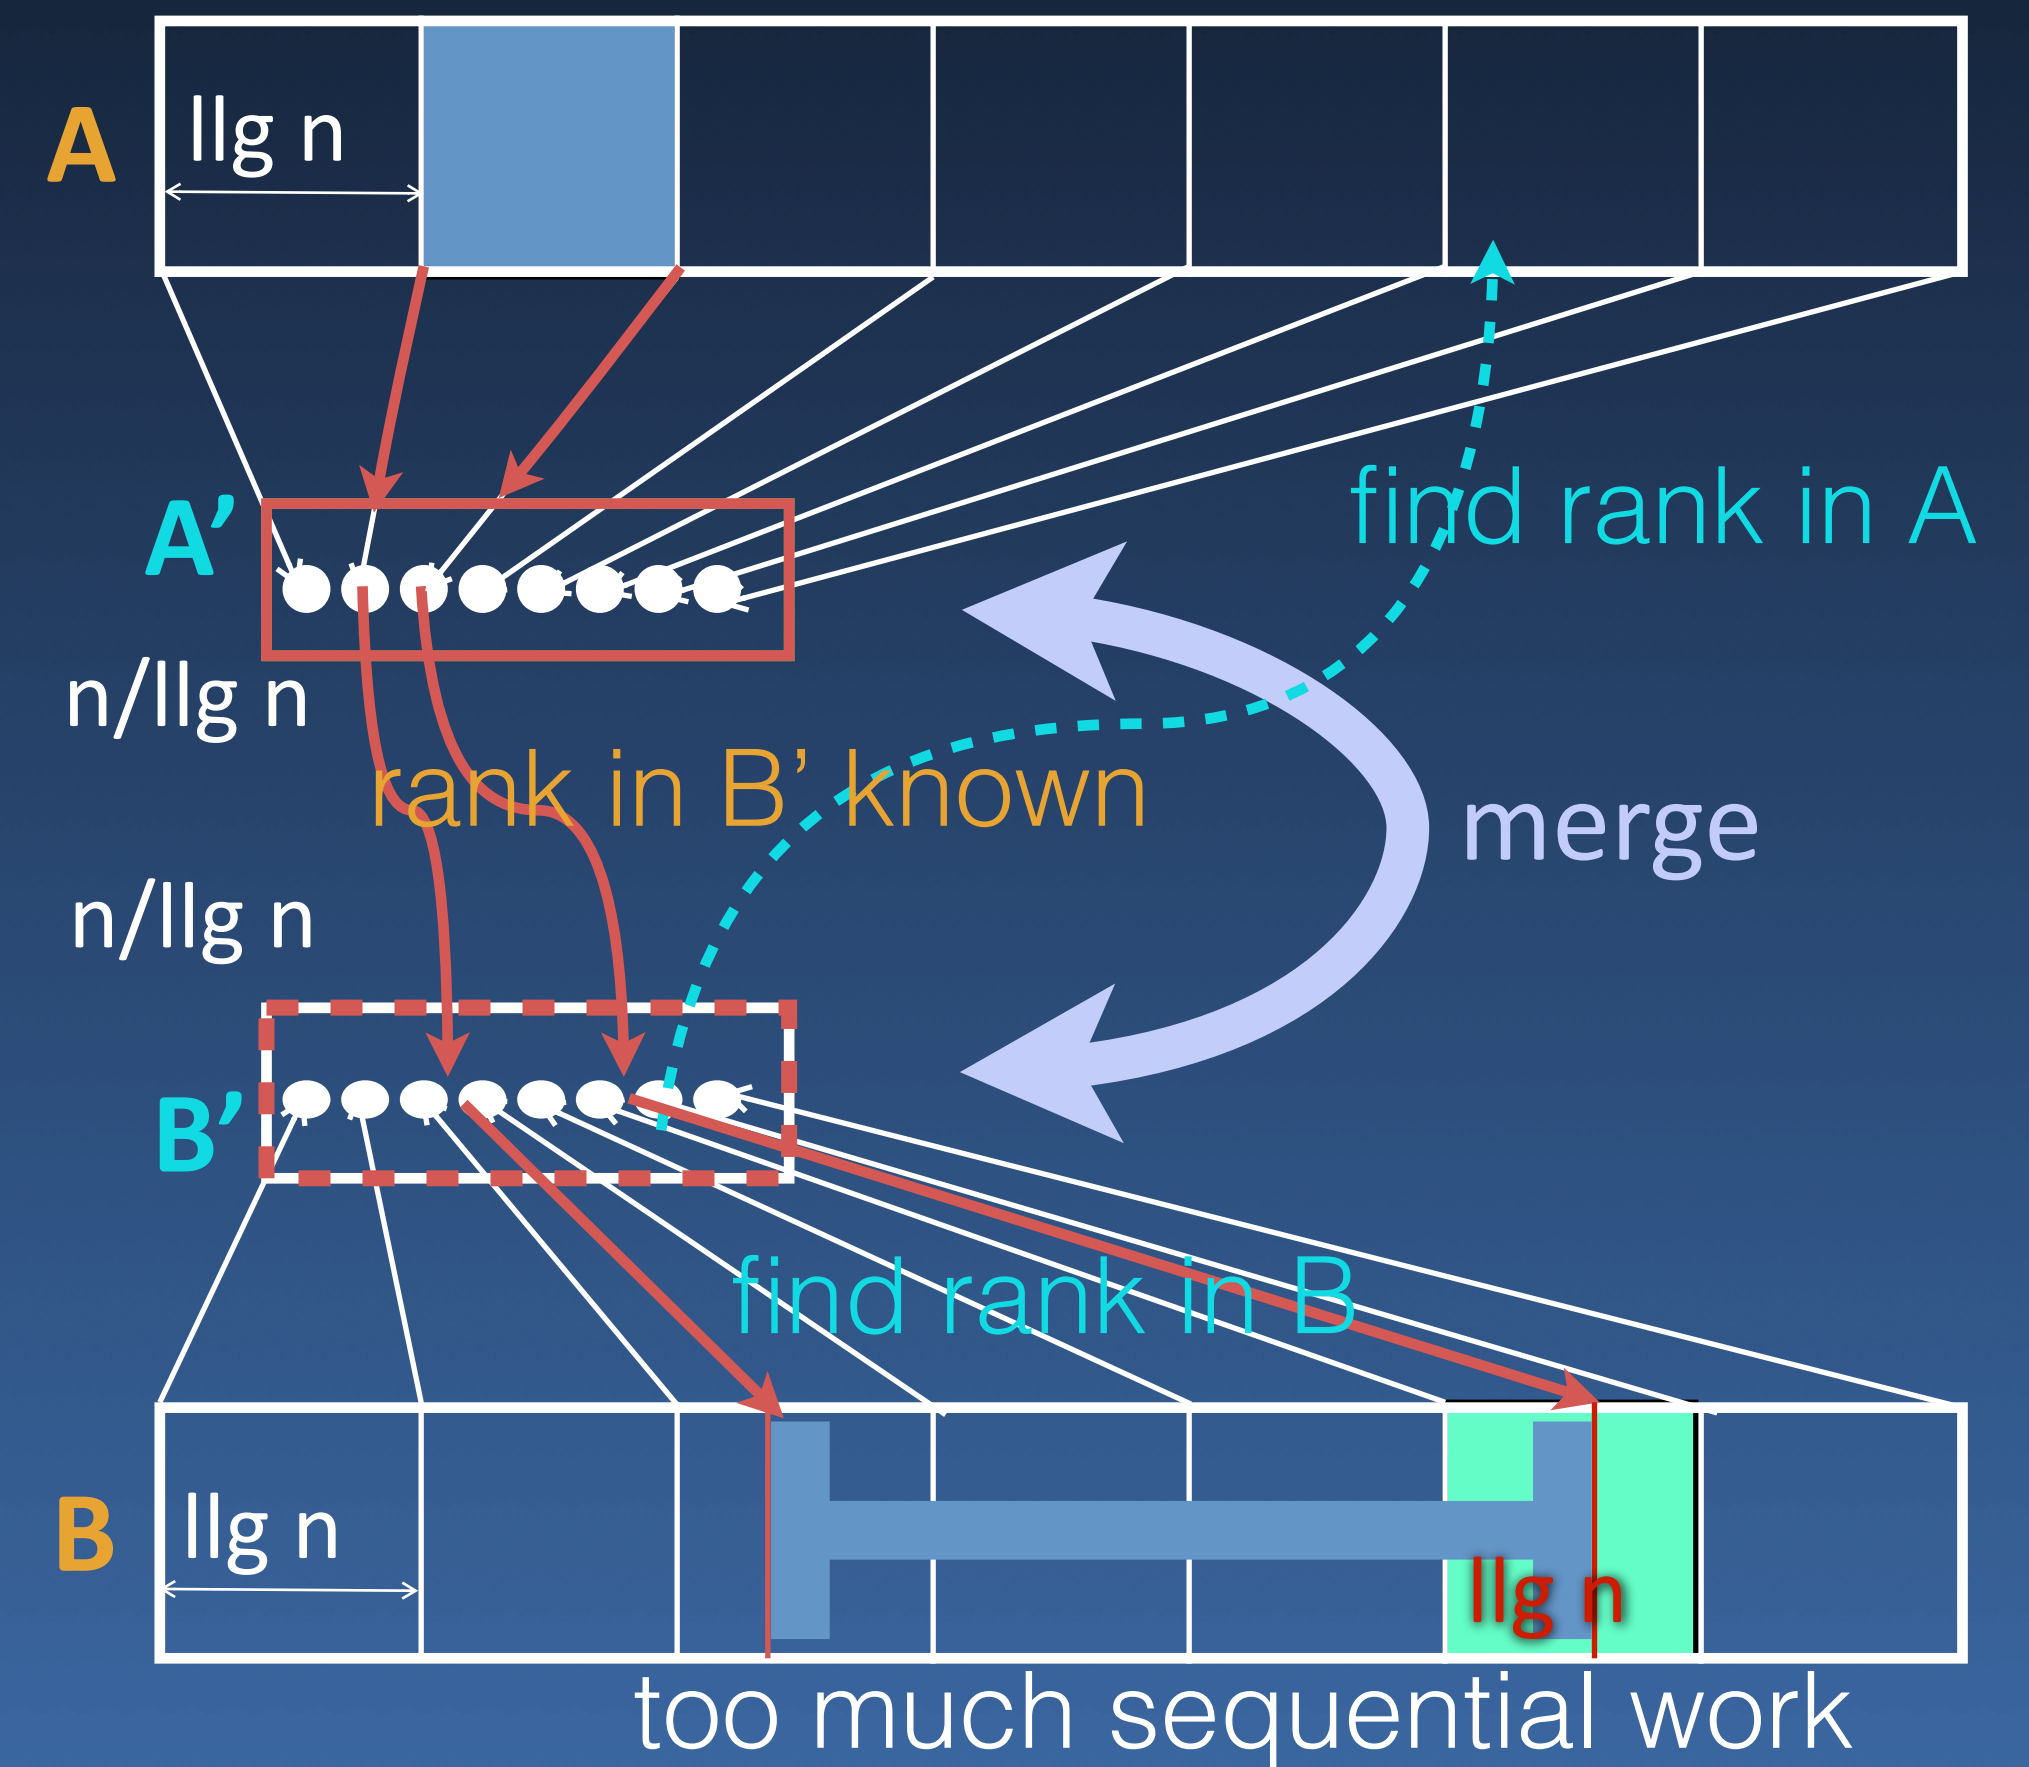
→ B_1, B_2, \dots

- Select first element of each block

→ $A' = p_1, p_2, \dots$

→ $B' = q_1, q_2, \dots$

- Now merge separated blocks of A and B



Optimal Merge (A,B)

- Use the fast, non-optimal algorithm on small enough subsets

- Subdivide A and B into blocks of size $\lg n$ ($\lg = \log \log$)

→ A_1, A_2, \dots

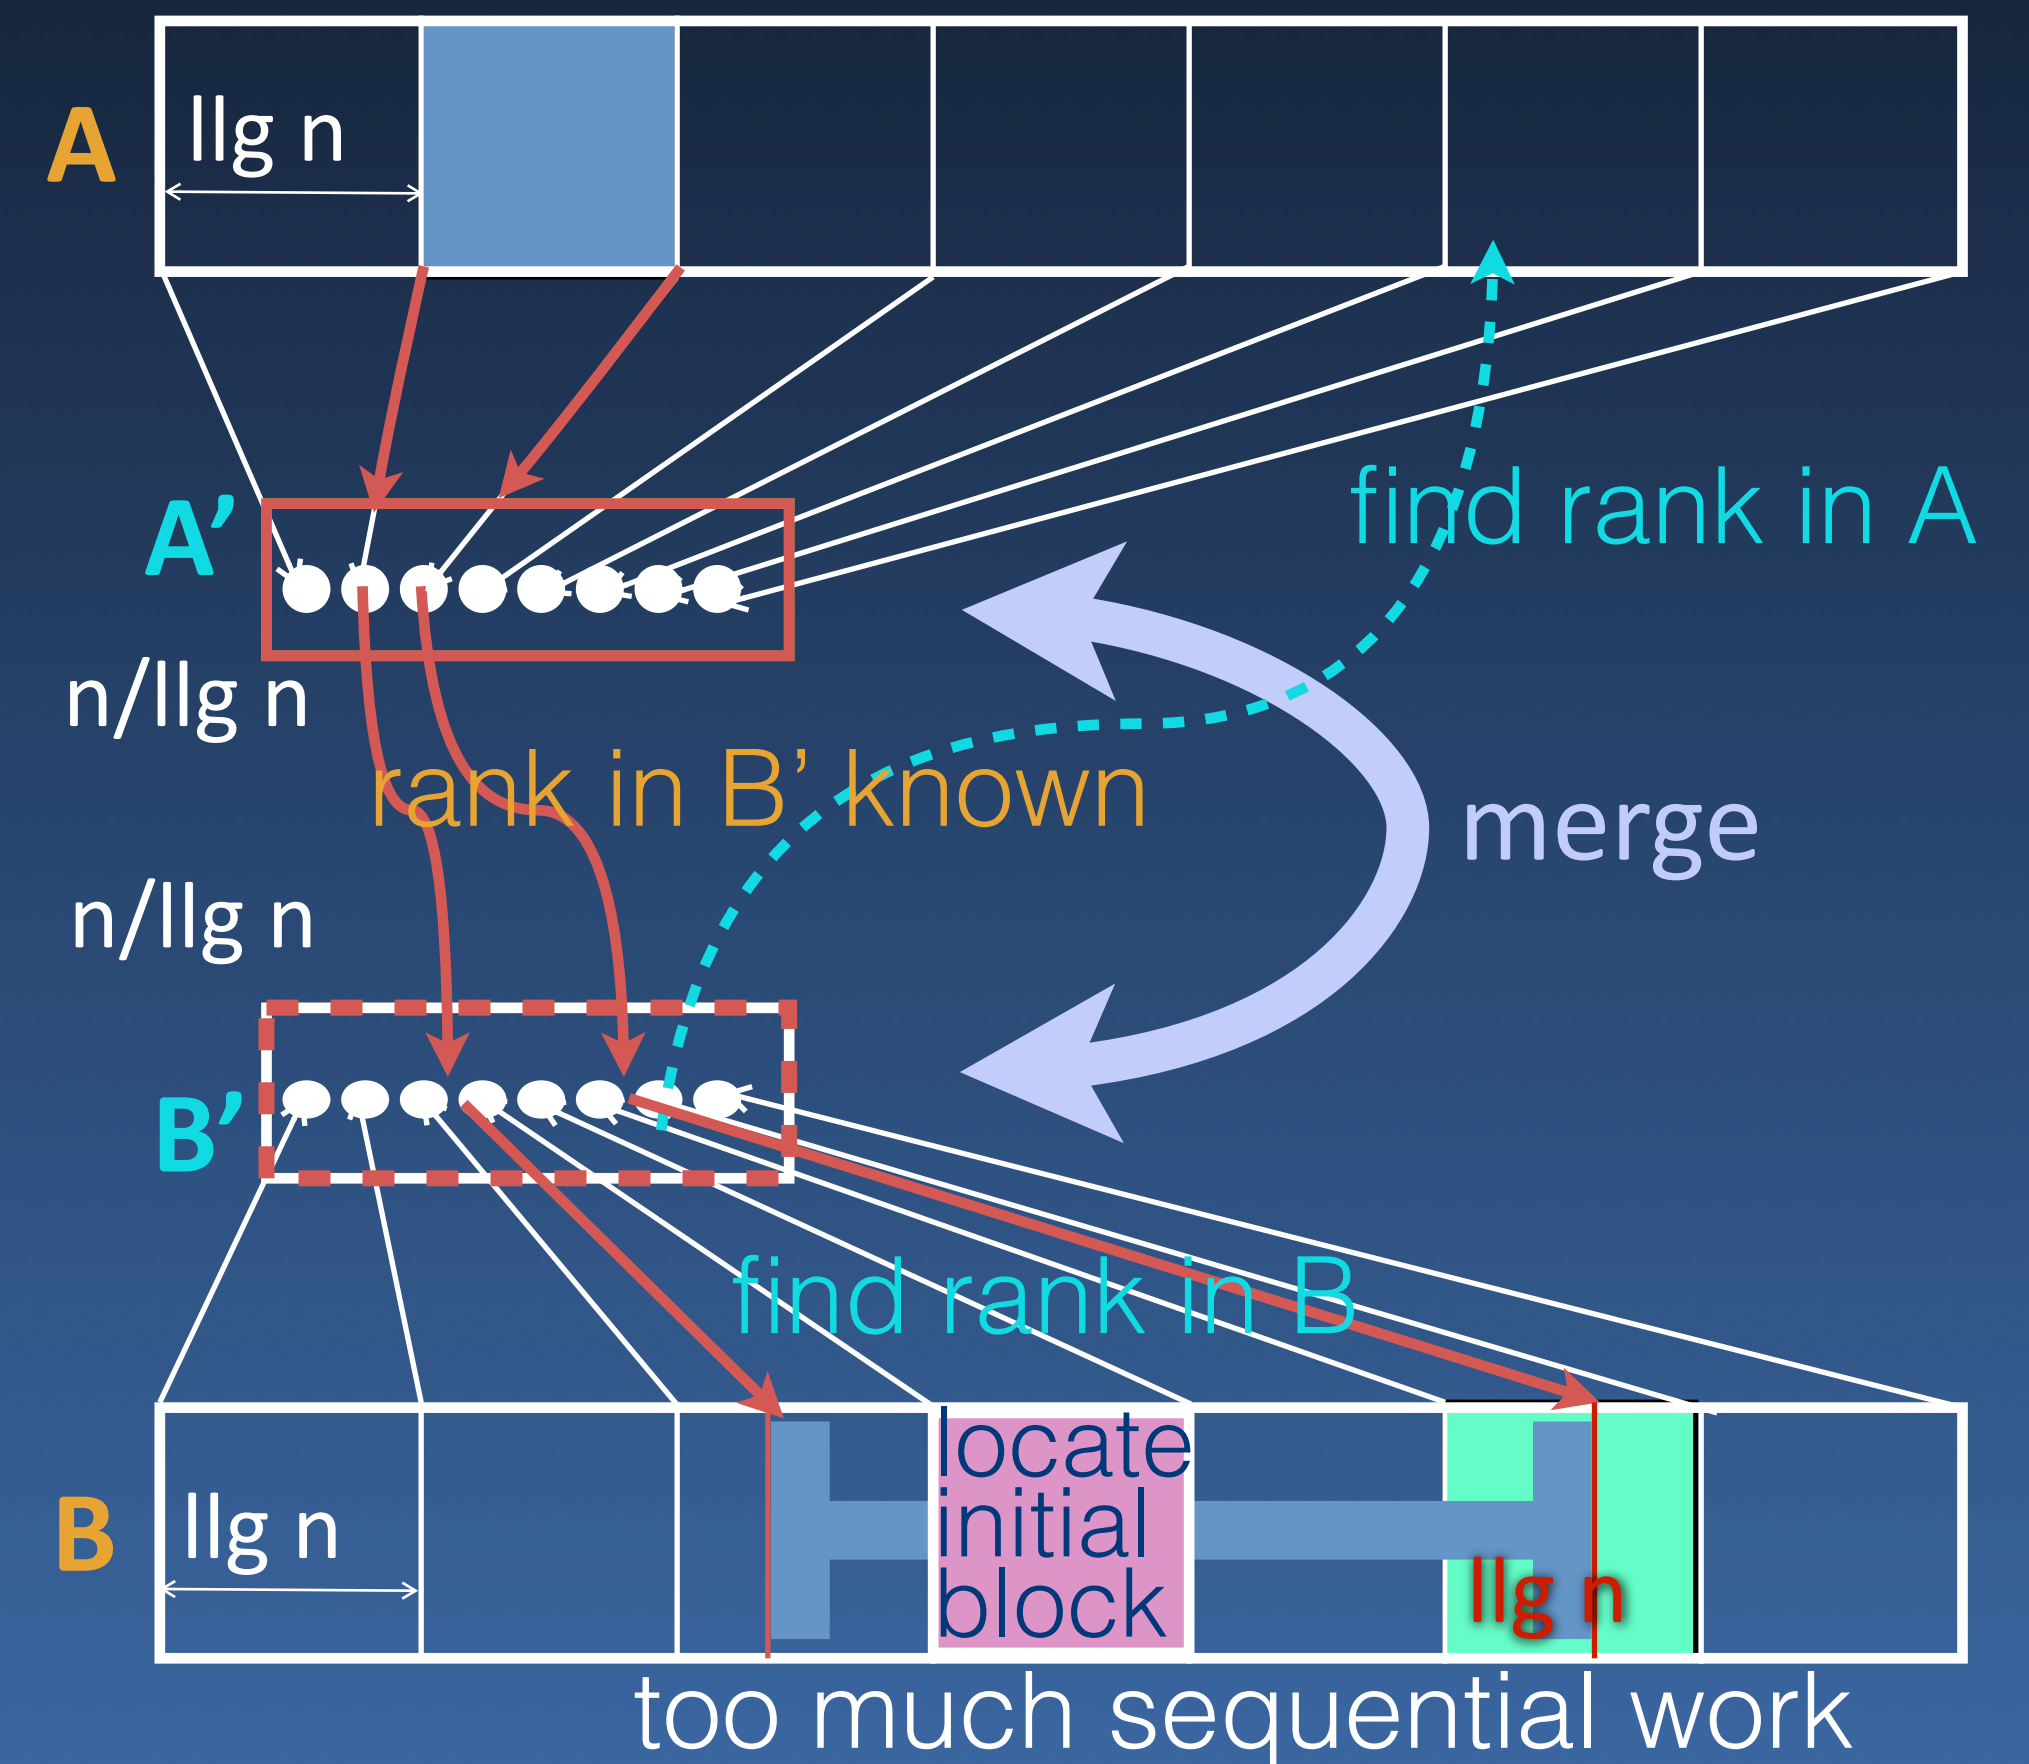
→ B_1, B_2, \dots

- Select first element of each block

→ $A' = p_1, p_2, \dots$

→ $B' = q_1, q_2, \dots$

- Now merge separated blocks of A and B



Optimal Merge (A,B)

- Use the fast, non-optimal algorithm on small enough subsets

- Subdivide A and B into blocks of size $\lg n$ ($\lg = \log \log$)

→ A_1, A_2, \dots

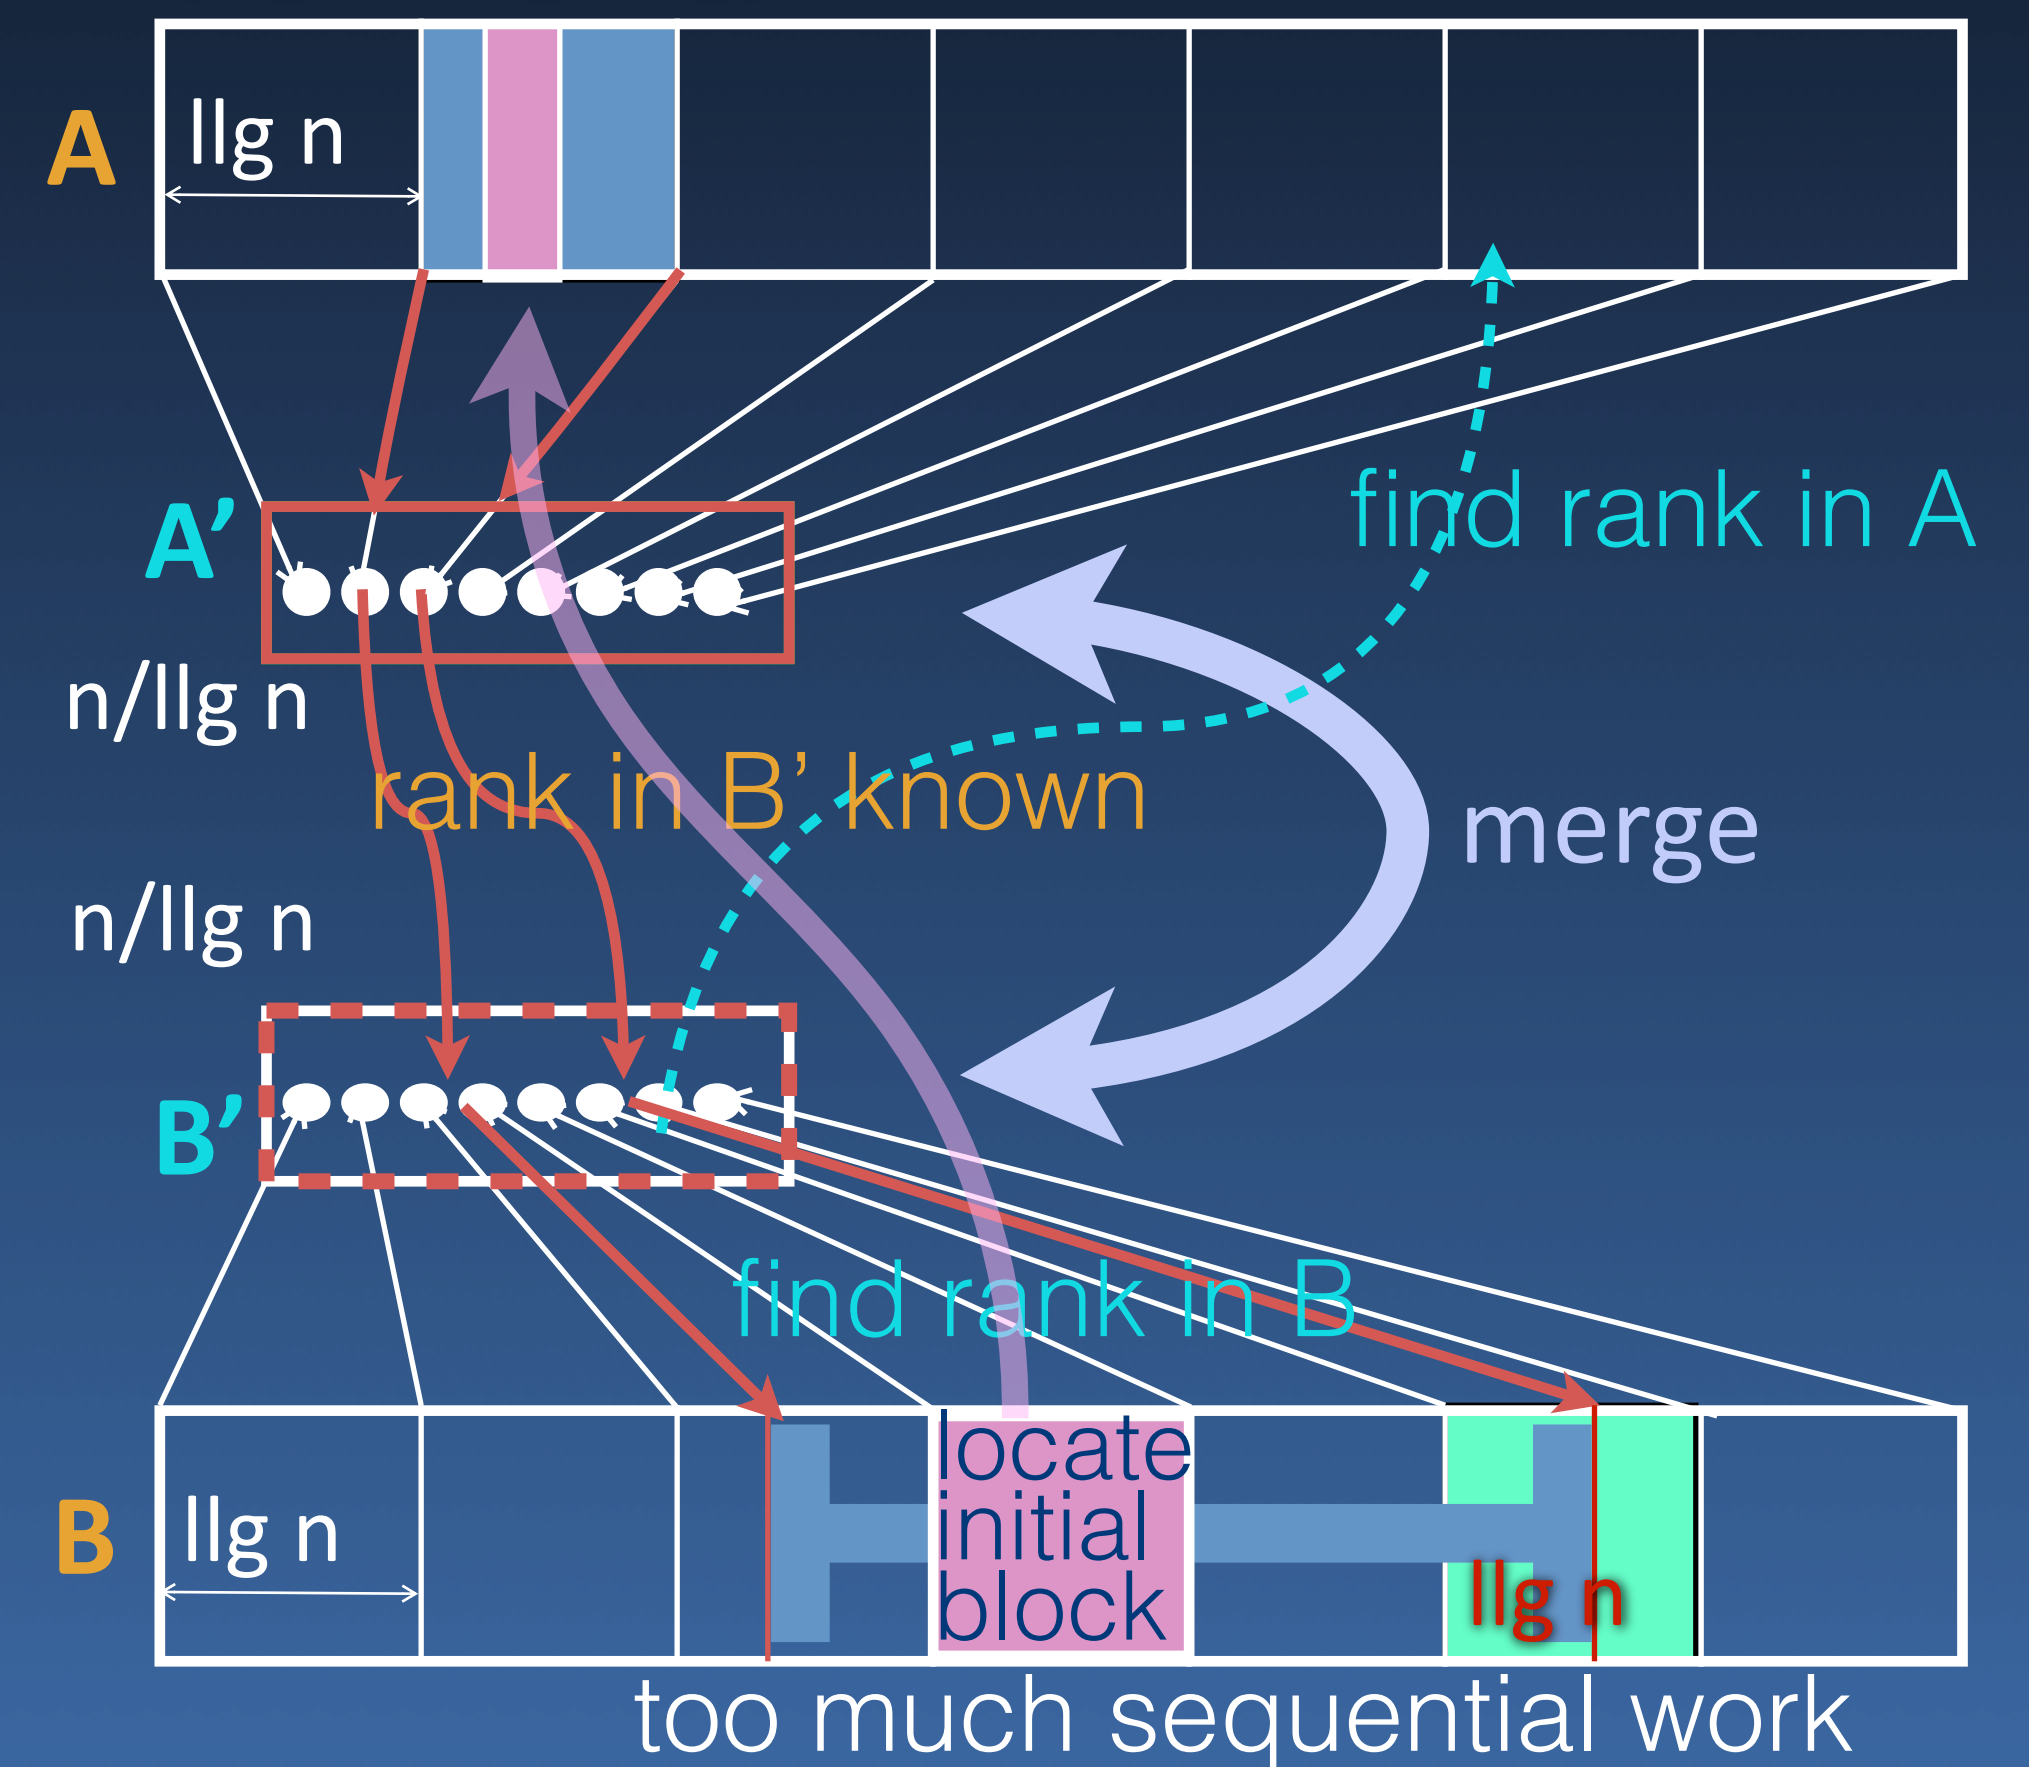
→ B_1, B_2, \dots

- Select first element of each block

→ $A' = p_1, p_2, \dots$

→ $B' = q_1, q_2, \dots$

- Now merge separated blocks of A and B



Optimal Merge (A,B)

- Use the fast, non-optimal algorithm on small enough subsets

- Subdivide A and B into blocks of size $\lg n$ ($\lg = \log \log$)

→ A_1, A_2, \dots


→ B_1, B_2, \dots

- Select first element of each block

→ $A' = p_1, p_2, \dots$

→ $B' = q_1, q_2, \dots$

- Merge $O(\lg n)$ sized blocks $O(n/\lg n)$ times



Optimal Merge (A,B) (Analysis)

1. Merge A' and B' – find $\text{Rank}(A':B')$, $\text{Rank}(B':A')$

→ Use fast non-optimal algorithm

- Time = $O(\log \log n)$, Work = $O(n)$


2. Compute $\text{Rank}(A':B)$ and $\text{Rank}(B':A)$

→ If $\text{Rank}(p_i, B)$ is r_i , p_i lies in block B_{r_i}

- Sequentially: Time = $O(\log \log n)$, Work = $O(n)$

3. Compute ranks of remaining elements

- Sequentially: Time = $O(\log \log n)$, Work = $O(n)$




Min-find

Input: array A with n elements

Algorithm A1 using $O(n^2)$ processors:

CRCW



Min-find


Input: array A with n elements

Algorithm A1 using $O(n^2)$ processors:

forall i in [0:n)

M[i] = 1

CRCW



Min-find


Input: array A with n elements

Algorithm A1 using $O(n^2)$ processors:

forall i in [0:n)

M[i] = 1

CRCW



Min-find

Input: array A with n elements

Algorithm A1 using $O(n^2)$ processors:

forall i in [0:n)


$M[i] = 1$

forall i,j in [0:n)

if $i \neq j \ \&\& \ A[i] < A[j]$

$M[j] = 0$

CRCW



Min-find

Input: array A with n elements

Algorithm A1 using $O(n^2)$ processors:

forall i in [0:n)

M[i] = 1

forall i,j in [0:n)

if $i \neq j$ && $A[i] < A[j]$


M[j] = 0

forall i in (0:n]

if M[i]=1

min = A[i]


CRCW



$O(1)$ time, $O(n^2)$ work: Not optimal

- **Balanced Binary tree**
 - $O(\log n)$ time
 - $O(n)$ work => **Optimal**
- **Make the tree branch quicker**
 - Number of children of node $u = \sqrt{n_u}$
 - ▶ if the number of leaves in u 's subtree is n_u
 - Works well if the operation at each node is fast, say, $O(1)$
- **Use Accelerated cascading**

From n^2 processors to $n\sqrt{n}$



Algorithm A2

Step 1: Partition into disjoint blocks of size \sqrt{n}

From n^2 processors to $n\sqrt{n}$

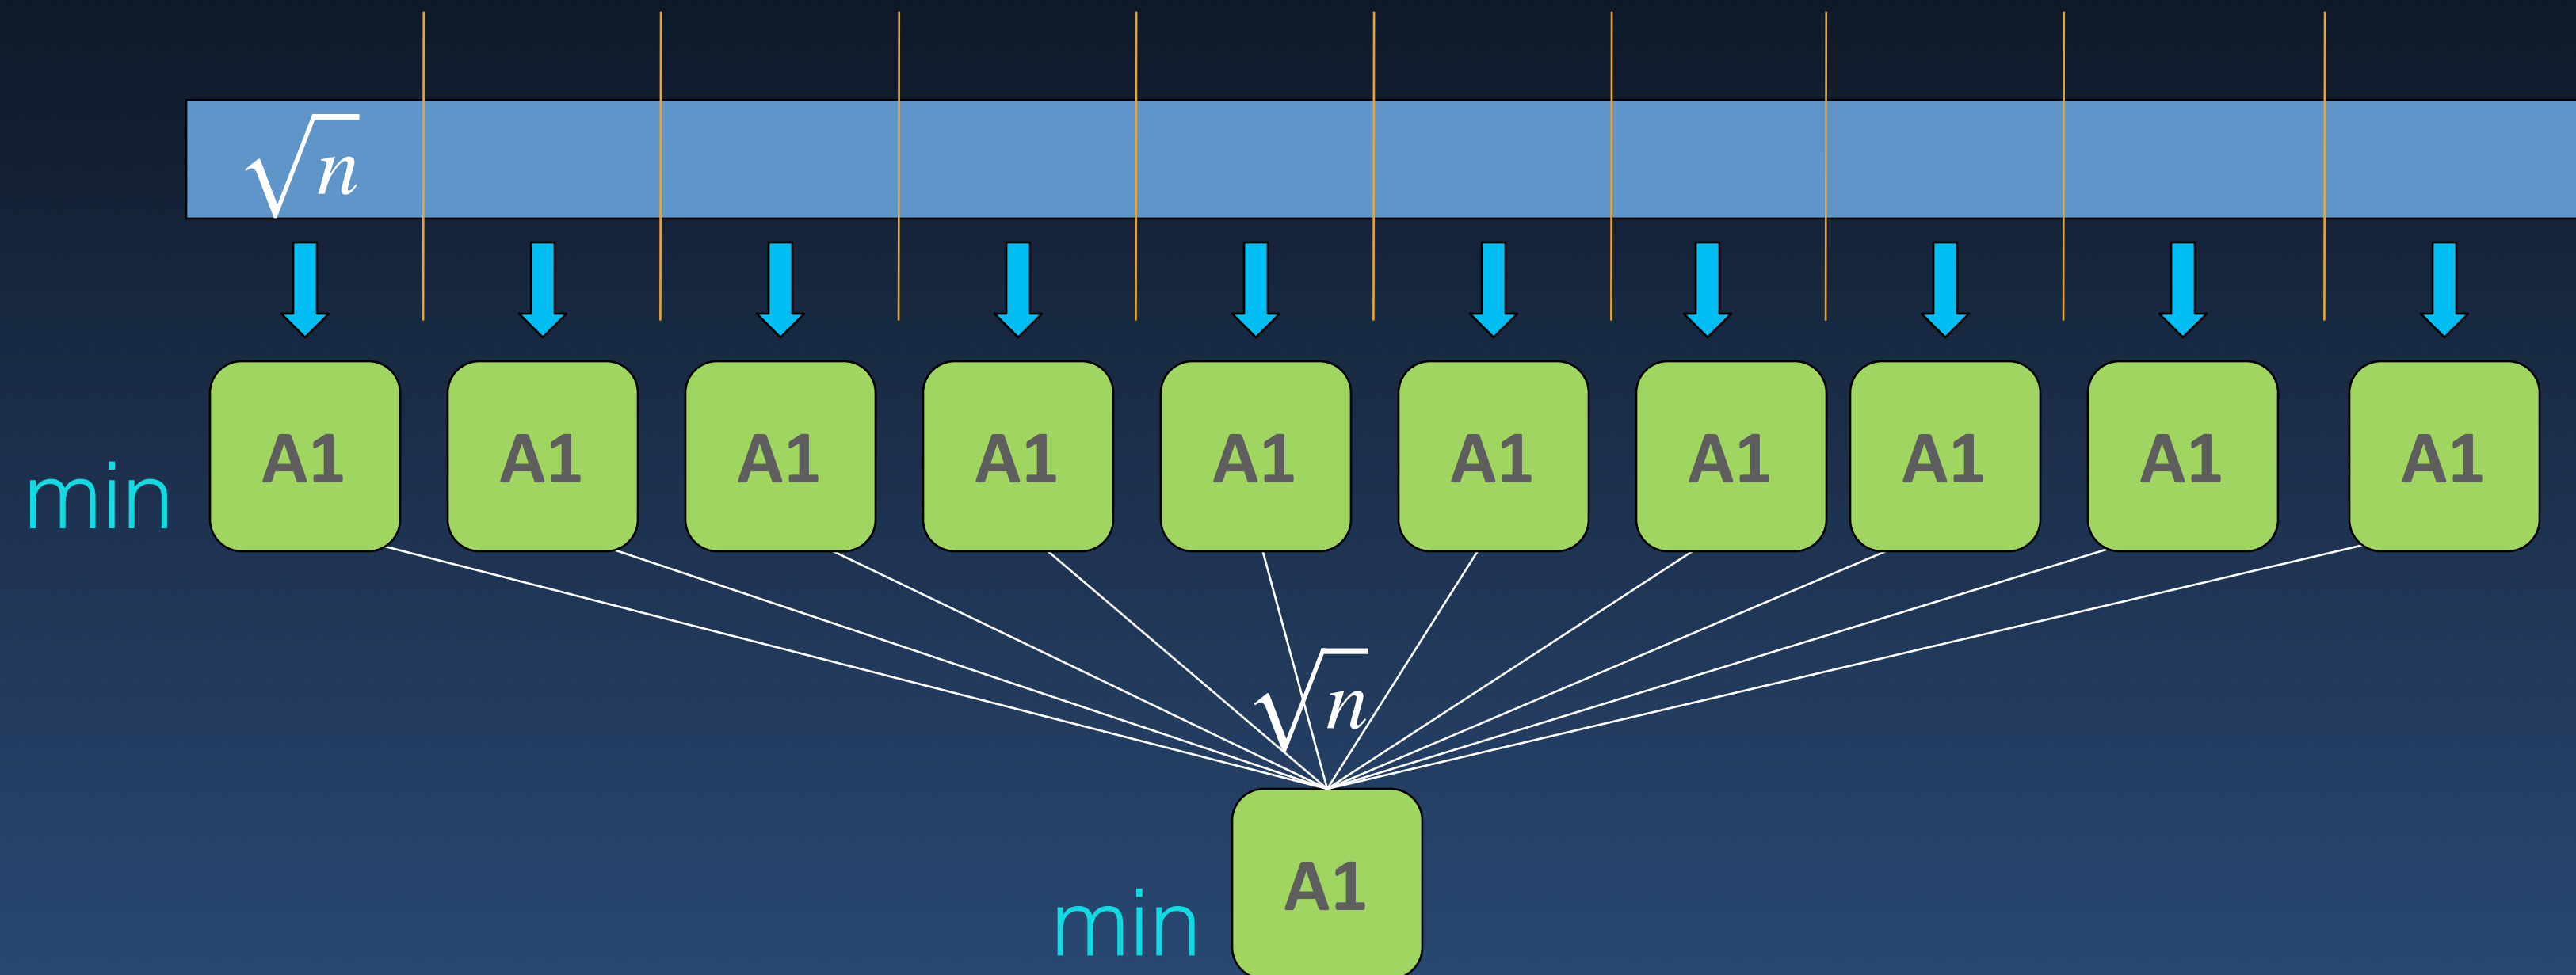


Algorithm A2

Step 1: Partition into disjoint blocks of size \sqrt{n}

Step 2: Apply A1 to each block

From n^2 processors to $n\sqrt{n}$




Algorithm A2

Step 1: Partition into disjoint blocks of size \sqrt{n}

Step 2: Apply A1 to each block

Step 3: Apply A1 to the results from the step 2

From n^2 processors to $n\sqrt{n}$



Algorithm A2

Step 1: Partition into disjoint blocks of size \sqrt{n}

Step 2: Apply A1 to each block

Step 3: Apply A1 to the results from the step 2

A2 work

$$n\sqrt{n}$$

$$n$$

From $n\sqrt{n}$ processors to $n^{1+1/4}$



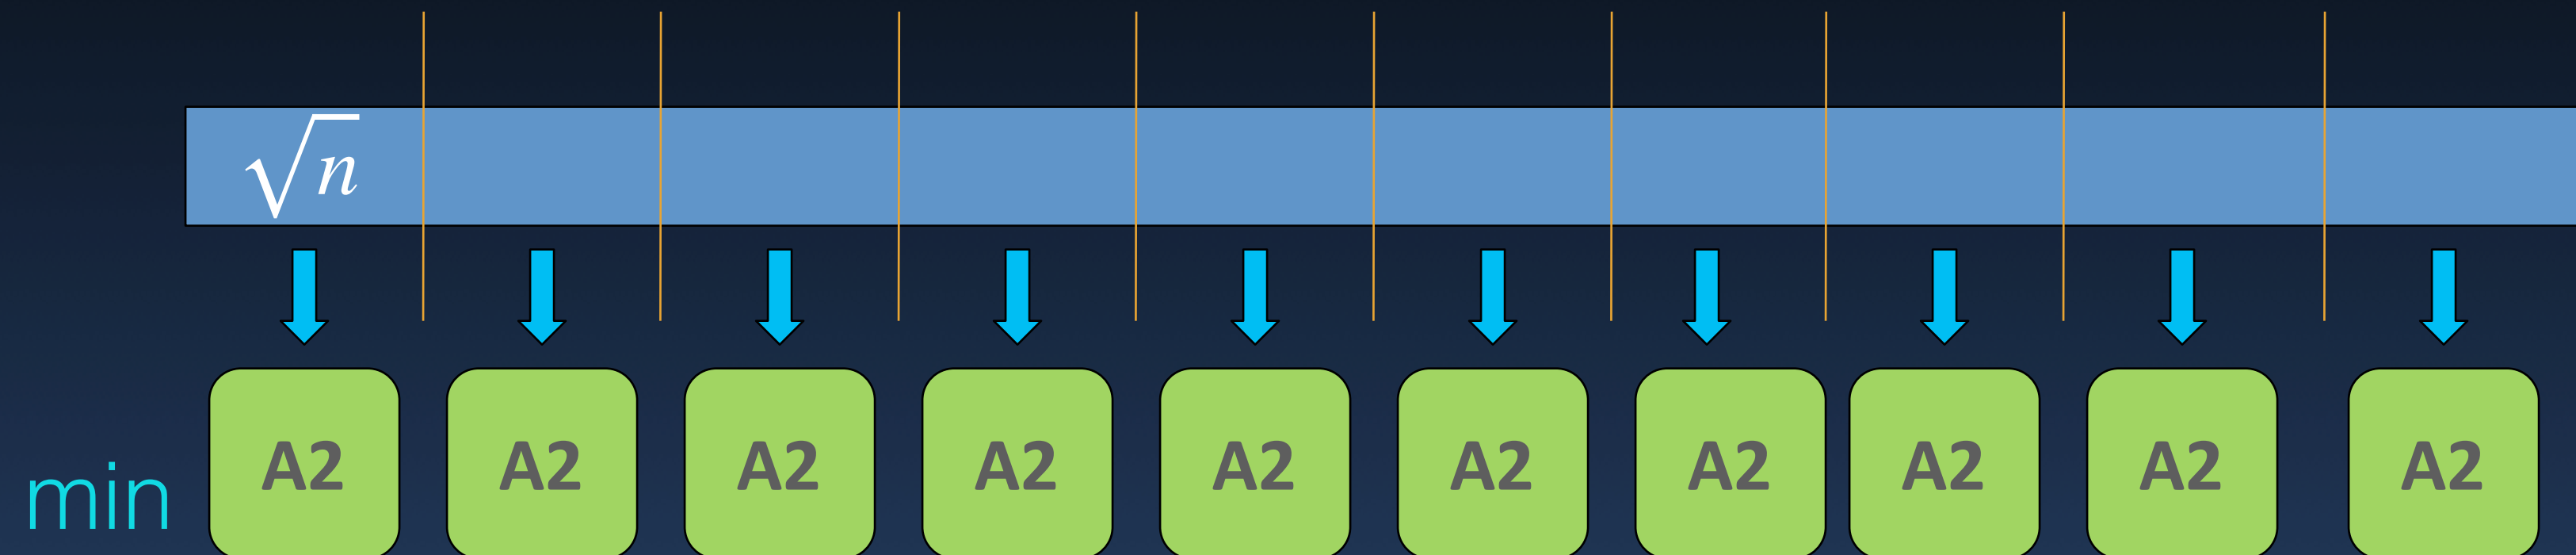
Algorithm A3

- Step 1: Partition into disjoint blocks of size
- Step 2: Apply A2 to each block
- Step 3: Apply A2 to the results from the step 2

A2 work


$n\sqrt{n}$
 n

From $n\sqrt{n}$ processors to $n^{1+1/4}$

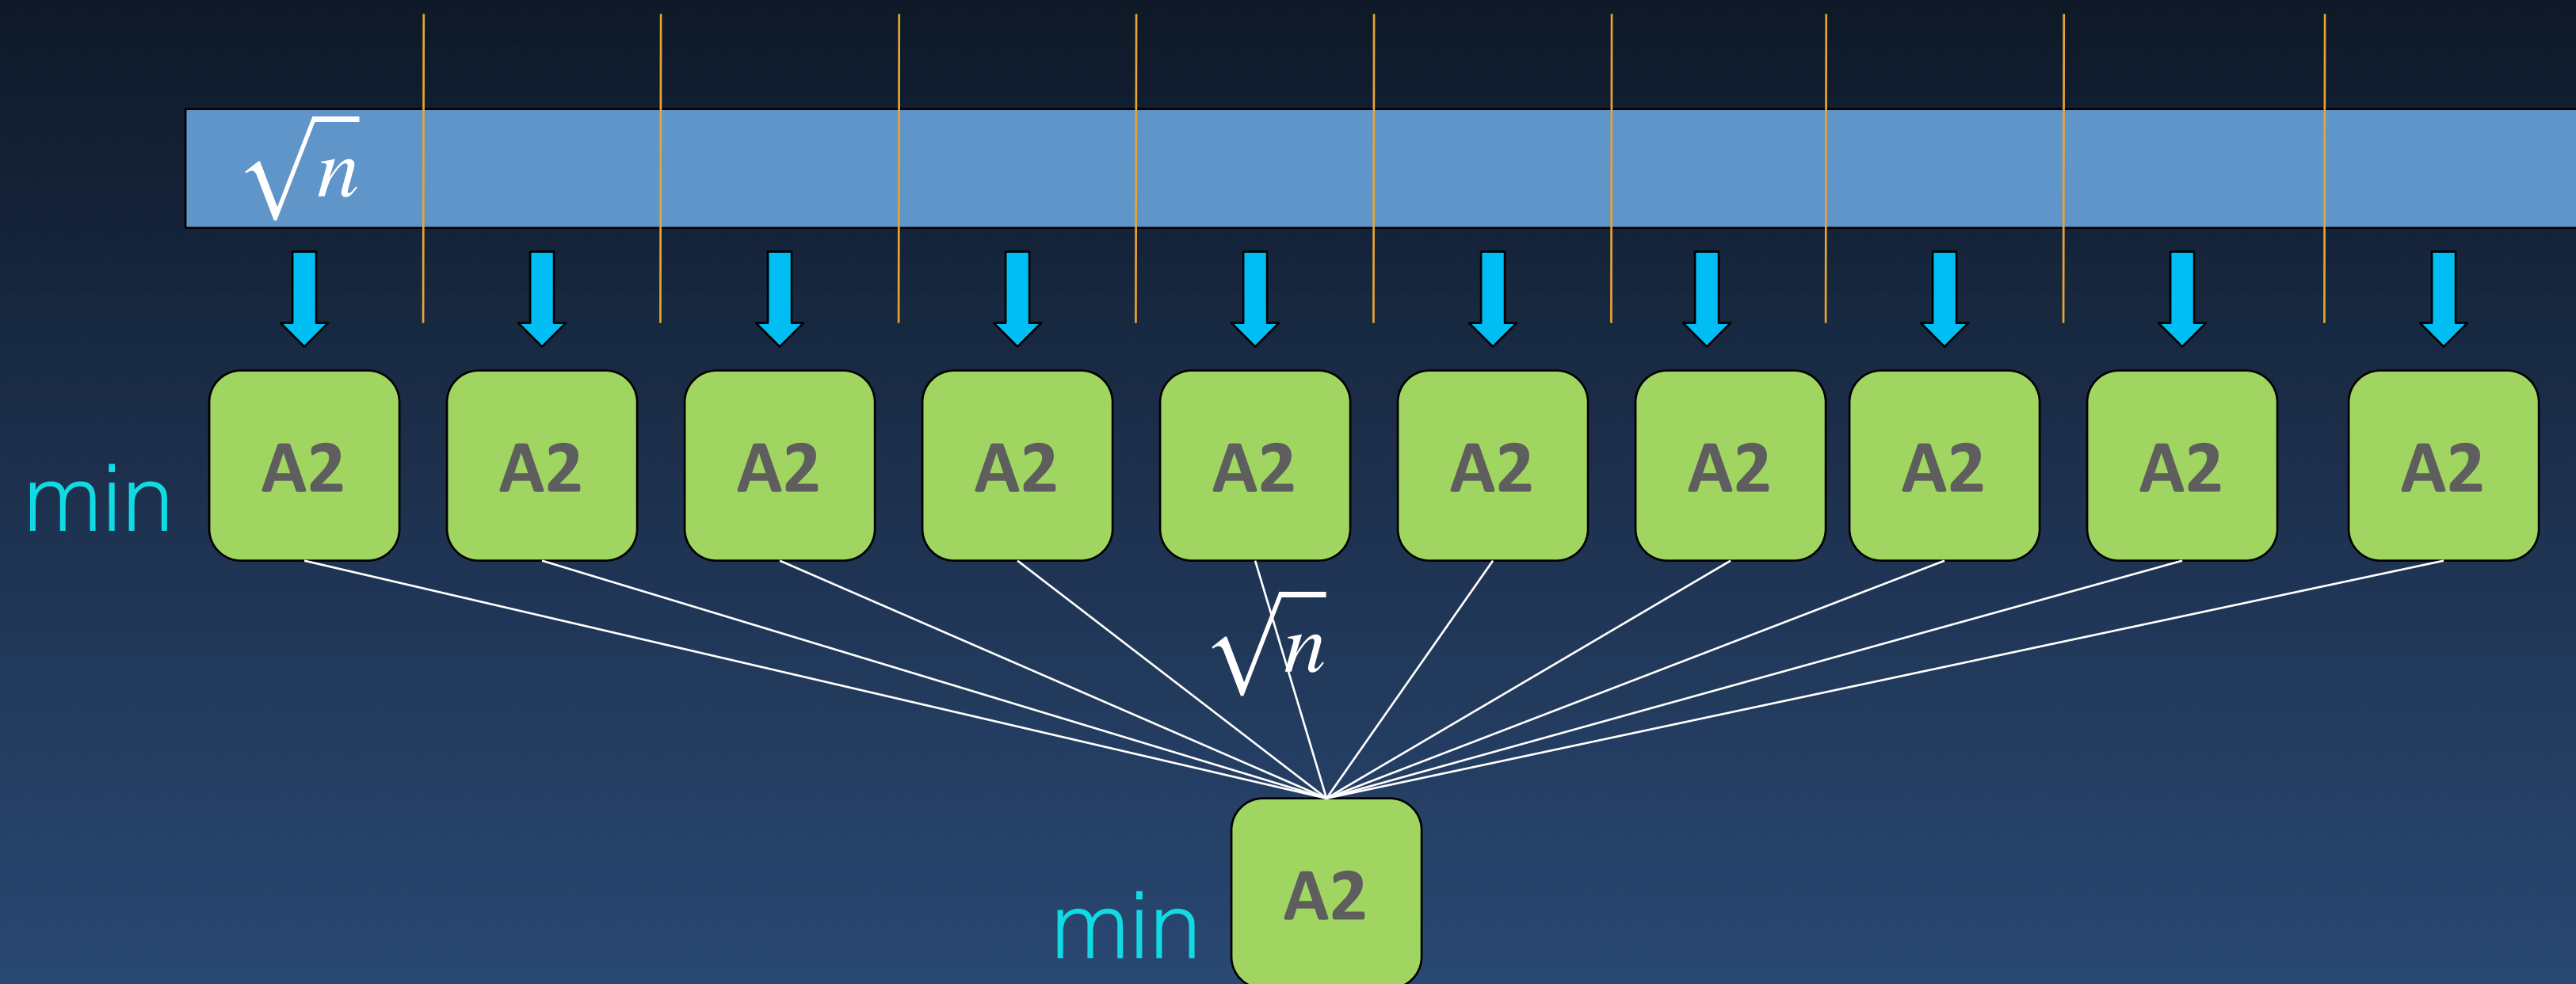


Algorithm A3

- Step 1: Partition into disjoint blocks of size
- Step 2: Apply A2 to each block
- Step 3: Apply A2 to the results from the step 2



From $n\sqrt{n}$ processors to $n^{1+1/4}$




Algorithm A3

- Step 1: Partition into disjoint blocks of size
- Step 2: Apply A2 to each block
- Step 3: Apply A2 to the results from the step 2

A2 work

$$\frac{n\sqrt{n}}{n}$$

From $n\sqrt{n}$ processors to $n^{1+1/4}$



Algorithm A3

- Step 1: Partition into disjoint blocks of size
- Step 2: Apply A2 to each block
- Step 3: Apply A2 to the results from the step 2

A3 work

$$n^{\frac{1}{2}}n^{\frac{3}{4}}$$
$$n^{\frac{3}{4}}$$

A2 work

$$n\sqrt{n}$$
$$n$$

Algorithm A_{k+1}

1. Partition input array C (size n) into disjoint blocks of size $n^{1/2}$ each
2. Solve for each block in parallel using algorithm A_k
3. Re-apply A_k to the results of step 2: minimum of $n^{1/2}$ minima

Algorithm A_{k+1}

1. Partition input array C (size n) into disjoint blocks of size $n^{1/2}$ each
2. Solve for each block in parallel using algorithm A_k
3. Re-apply A_k to the results of step 2: minimum of $n^{1/2}$ minima

A_1		A_2		A_3		..				
n^2	\rightarrow	$n^{1+1/2}$	\rightarrow	$n^{1+1/4}$	\rightarrow	$n^{1+1/8}$	\rightarrow	$n^{1+1/2^k}$..	$\sim n^{1+\epsilon}$

Algorithm A_∞ takes ?? with $n^{1+\epsilon}$ processors

Doubly-log depth tree: $n^{\frac{1}{2^i}} = O(1)$ at leaf

Algorithm A_{k+1}

1. Partition input array C (size n) into disjoint blocks of size $n^{1/2}$ each
2. Solve for each block in parallel using algorithm A_k
3. Re-apply A_k to the results of step 2: minimum of $n^{1/2}$ minima

A_1

A_2


A_3

..

$$n^2 \rightarrow n^{1+1/2} \rightarrow n^{1+1/4} \rightarrow n^{1+1/8} \rightarrow n^{1+1/2^k} \dots \sim n^{1+\epsilon}$$


Algorithm A_∞ takes ?? with $n^{1+\epsilon}$ processors

Doubly-log depth tree: $n^{\frac{1}{2^i}} = O(1)$ at leaf



Algorithm A_{k+1}

1. Partition input array C (size n) into disjoint blocks of size $n^{1/2}$ each
2. Solve for each block in parallel using algorithm A_k
3. Re-apply A_k to the results of step 2: minimum of $n^{1/2}$ minima



A_1	A_2	A_3	..
n^2	$\rightarrow n^{1+1/2}$	$\rightarrow n^{1+1/4}$	$\rightarrow n^{1+1/8} \rightarrow n^{1+1/2^k} \dots \sim n^{1+\epsilon}$

Algorithm A_∞ takes ?? with $n^{1+\epsilon}$ processors


Doubly-log depth tree: $n^{\frac{1}{2^i}} = O(1)$ at leaf

$n \log \log n$ work, $\log \log n$ time

- Constant-time algorithm
 - $O(n^2)$ work
- $O(\log n)$ Balanced tree approach
 - $O(n)$ work (Work-Optimal)
- $O(\log \log n)$ Doubly-log depth tree approach
 - $O(n \log \log n)$ work
 - Degree is high at the root, reduces going down
 - ▶ #Children of node $u = \sqrt{(\text{\#nodes in tree rooted at } u)}$
 - ▶ Depth = $O(\log \log n)$


Accelerated Cascading

- Solve recursively
- Start bottom-up with the optimal algorithm
 - until the problem sizes is smaller
- Switch to fast (non-optimal algorithm)
 - A few small problems solved fast but non-work-optimally
- Min Find:
 - Optimal algorithm for lower $\log \log n$ levels
 - Then switch to $O(n \log \log n)$ -work algorithm



Accelerated Cascading

- Solve recursively
- Start bottom-up with the optimal algorithm
 - until the problem sizes is smaller
- Switch to fast (non-optimal algorithm)
 - A few small problems solved fast but non-work-optimally
- Min Find:
 - Optimal algorithm for lower $\log \log n$ levels
 - Then switch to $O(n \log \log n)$ -work algorithm



$O(n)$ work, $O(\log \log n)$ time

- Choose the pivot
 - ➔ Select median?
- Subdivide into two groups
 - ➔ Group sizes linearly related with high probability (expect $\log(n)$ rounds)
- Sort each independently in parallel

Quick Sort

- Choose the pivot
 - ➔ Select median?
- Subdivide into two groups
 - ➔ Group sizes linearly related with high probability (expect $\log(n)$ rounds)
- Sort each independently in parallel

```
QuickSort(int A[], int first, int last)
{
    Select random m in [first:last]
    // A[m] is pivot
    forall i in [first:last]
        flag[i] = A[i] < A[m];
    Split(A); // Separate flag values 0 and 1
              // Prefix sum, A[P[m]] = A[m]
    Quicksort A[first:k-1] and A[k+1:last]
}
```

Quick Sort

- Choose the pivot
 - ➔ Select median?
- Subdivide into two groups
 - ➔ Group sizes linearly related with high probability (expect $\log(n)$ rounds)
- Sort each independently in parallel

```
QuickSort(int A[], int first, int last)
{
    Select random m in [first:last]
    // A[m] is pivot
    forall i in [first:last]
        flag[i] = A[i] < A[m];
    Split(A); // Separate flag values 0 and 1
              // Prefix sum, A[P[m]] = A[m]
    Quicksort A[first:k-1] and A[k+1:last]
}
```

	①	1	②	2	2	③	④
flag	1	0	1	0	0	1	1

Quick Sort

- Choose the pivot
 - ➔ Select median?
- Subdivide into two groups
 - ➔ Group sizes linearly related with high probability (expect $\log(n)$ rounds)
- Sort each independently in parallel

```
QuickSort(int A[], int first, int last)
{
    Select random m in [first:last]
    // A[m] is pivot
    forall i in [first:last]
        flag[i] = A[i] < A[m];
    Split(A); // Separate flag values 0 and 1
              // Prefix sum, A[P[m]] = A[m]
    Quicksort A[first:k-1] and A[k+1:last]
}
```

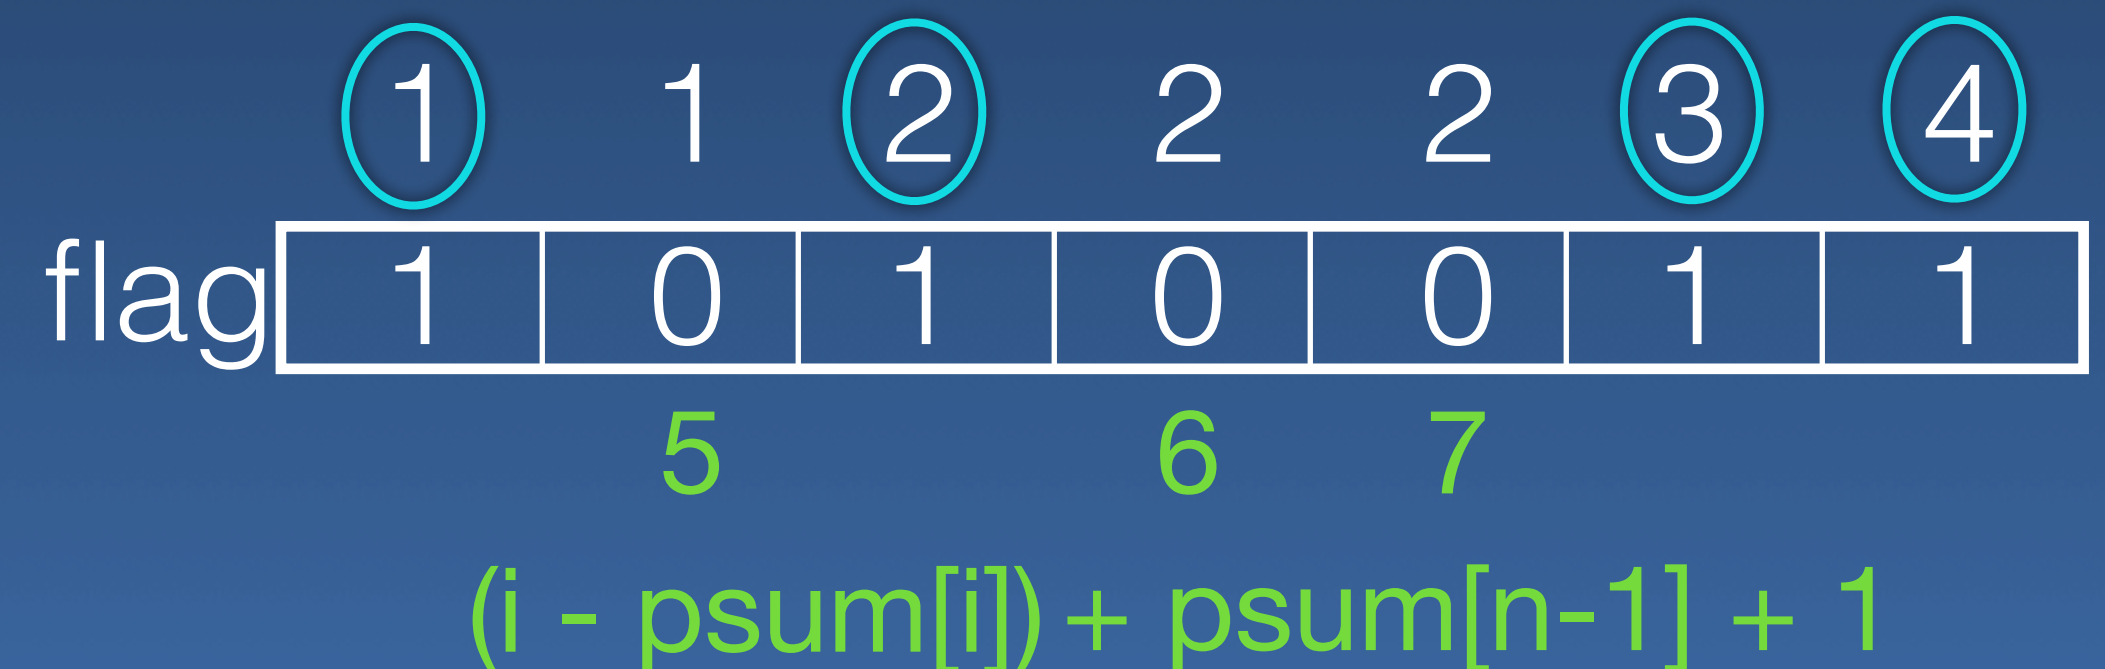
	①	1	②	2	2	③	④
flag	1	0	1	0	0	1	1

$(i - \text{psum}[i])$

Quick Sort

- Choose the pivot
 - ➔ Select median?
- Subdivide into two groups
 - ➔ Group sizes linearly related with high probability (expect $\log(n)$ rounds)
- Sort each independently in parallel

```
QuickSort(int A[], int first, int last)
{
    Select random m in [first:last]
    // A[m] is pivot
    forall i in [first:last]
        flag[i] = A[i] < A[m];
    Split(A); // Separate flag values 0 and 1
              // Prefix sum, A[P[m]] = A[m]
    Quicksort A[first:k-1] and A[k+1:last]
}
```



Quick Sort

- Choose the pivot
 - ➔ Select median?
- Subdivide into two groups
 - ➔ Group sizes linearly related with high probability (expect $\log(n)$ rounds)
- Sort each independently in parallel
- Time per round = $O(\log n)$
- Work per round = $O(n)$

```
QuickSort(int A[], int first, int last)
{
    Select random m in [first:last]
    // A[m] is pivot
    forall i in [first:last]
        flag[i] = A[i] < A[m];
    Split(A); // Separate flag values 0 and 1
              // Prefix sum, A[P[m]] = A[m]
    Quicksort A[first:k-1] and A[k+1:last]
}
```

	①	1	②	2	2	③	④
flag	1	0	1	0	0	1	1

$$(i - \text{psum}[i]) + \text{psum}[n-1] + 1$$

$$T(n) \sim T(n/2) + O(\log n)$$
$$W(n) \sim 2W(n/2) + O(n)$$

- Merge, Minima-find, Quicksort
- Combining non-work optimal fast algorithm with optimal less fast algorithm
 - ➔ Accelerated Cascading