

COL380

Introduction to
Parallel & Distributed Programming

- Formal Models of Parallel Computational
 - ➔ Actor model
 - ➔ PRAM model

- Simplify specifying, reasoning, analyzing algorithms
- Must abstract away many details
 - ➔ Should predict computability
- Should track performance
- General classes
 - ➔ Shared Memory vs Distributed Memory
 - ➔ Synchronous vs Asynchronous

- Actors
- Communicating Sequential Processes
- Parallel Random Access Machine
- Bulk Synchronous Parallel computation

- **Actors: autonomous computing agents**
 - ➔ No shared state; interact with each other using messages
 - ▶ Asynchronous, Lossless, Unordered
 - ▶ Have a (addressed) mailbox for communication (allows buffering)
- **Actors process messages in their mailbox, in response:**
 - ➔ sends zero or more messages
 - ➔ creates zero or more new actors
 - ➔ changes its own local state (impacts the next message, 'local computation')

[See: Hewitt, Bishop, Steiger, "A Universal Modular Actor Formalism for Artificial Intelligence," IJCAI 1973]


- Compose sequential processes passing messages
 - ➔ Synchronous: send completes when message received (and vice versa)
 - ➔ Processes names known to senders (used for send and receive)
- $\text{Sender?gotvalue} \parallel \text{Recipient!somevalue}$
- $\text{Guard; sender?P} \rightarrow \{\text{post arrival code}\}$
 - ➔ Wait for predicate to become true (or fail if it becomes false)

```
[ Sender1? msg -> process(msg, Q) ||  
  Sender2? msg -> process(msg, R) ||  
  Sender3? msg -> process(msg, S)  ]
```

(Selection)


[See Hoare, "Communicating sequential processes," CACM 21 (8), 1978]

PRAM Model



- Synchronous, Shared Mem
 - ➔ *Arbitrary* number of cells
- Arbitrary number of processors, each:
 - ➔ has local memory (*Arbitrary* number of cells)
 - ➔ knows its ID
 - ➔ can access a shared memory location in *constant* time

PRAM Model

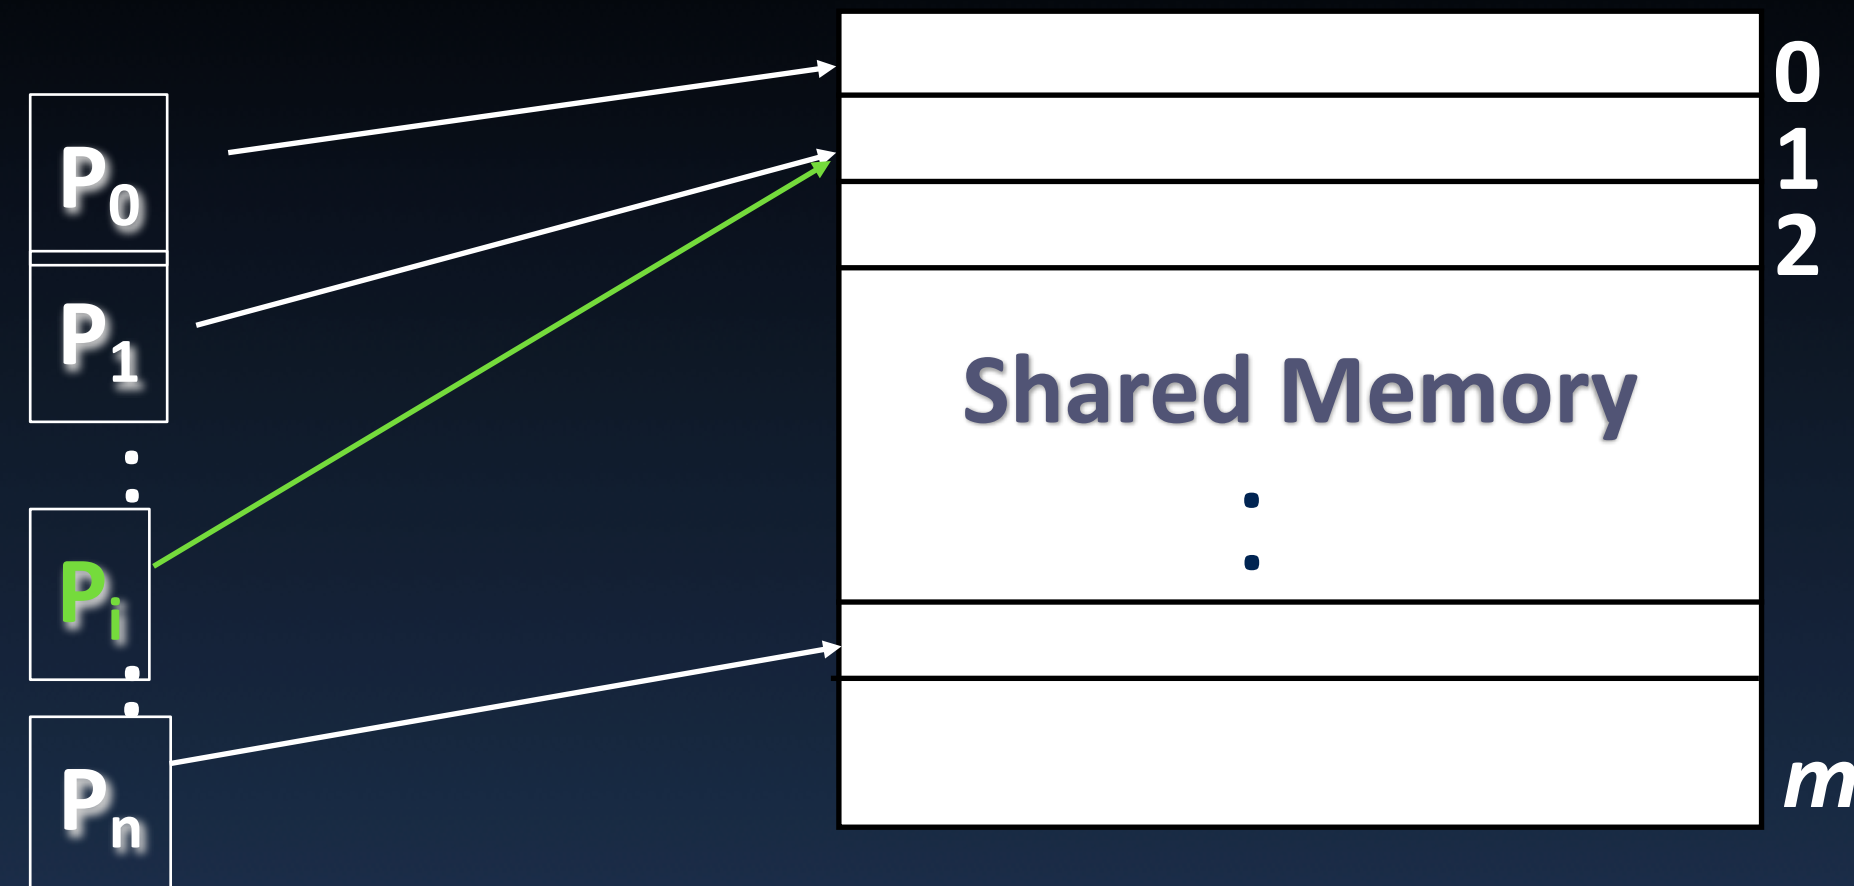


- Synchronous, Shared Mem
 - ➔ *Arbitrary* number of cells
- Arbitrary number of processors, each:
 - ➔ has local memory (*Arbitrary* number of cells)
 - ➔ knows its ID
 - ➔ can access a shared memory location in *constant* time

Unrealistic?

Can be often simulated

PRAM Model Steps



- At each time-step each P_i can:
 1. read some memory cell
 2. perform a local computation step
 3. write a memory cell (Read and write are in two phases)
 - ➔ Co-access may be restricted
- Thus, a pair of processors P_i and P_j can communicate in two steps
 - ➔ constant time

- Inputs/Outputs are placed at designated addresses
 - ➔ Technically also a 'start' protocol to activate processors
- Each instruction takes $O(1)$ time
- Processors are synchronous
 - ➔ Asynchronous PRAM models exist as well
- Cost analysis:
 - ➔ Cost, Work, Time (taken by the longest running processor)
 - ➔ Maximum number of active processors and memory cells

- EREW (Exclusive Read Exclusive Write)
 - ➔ Only one processors may read or write any given location in a step
- CREW (Concurrent Read Exclusive Write)
 - ➔ Many processors can simultaneously read a location, but only one may write
- CRCW (Concurrent Read Concurrent Write)
 - ➔ Many processors can read/write the same memory location
- ERCW (Exclusive Read Concurrent Write)
 - ➔ Not commonly used

Concurrent Write (CW)

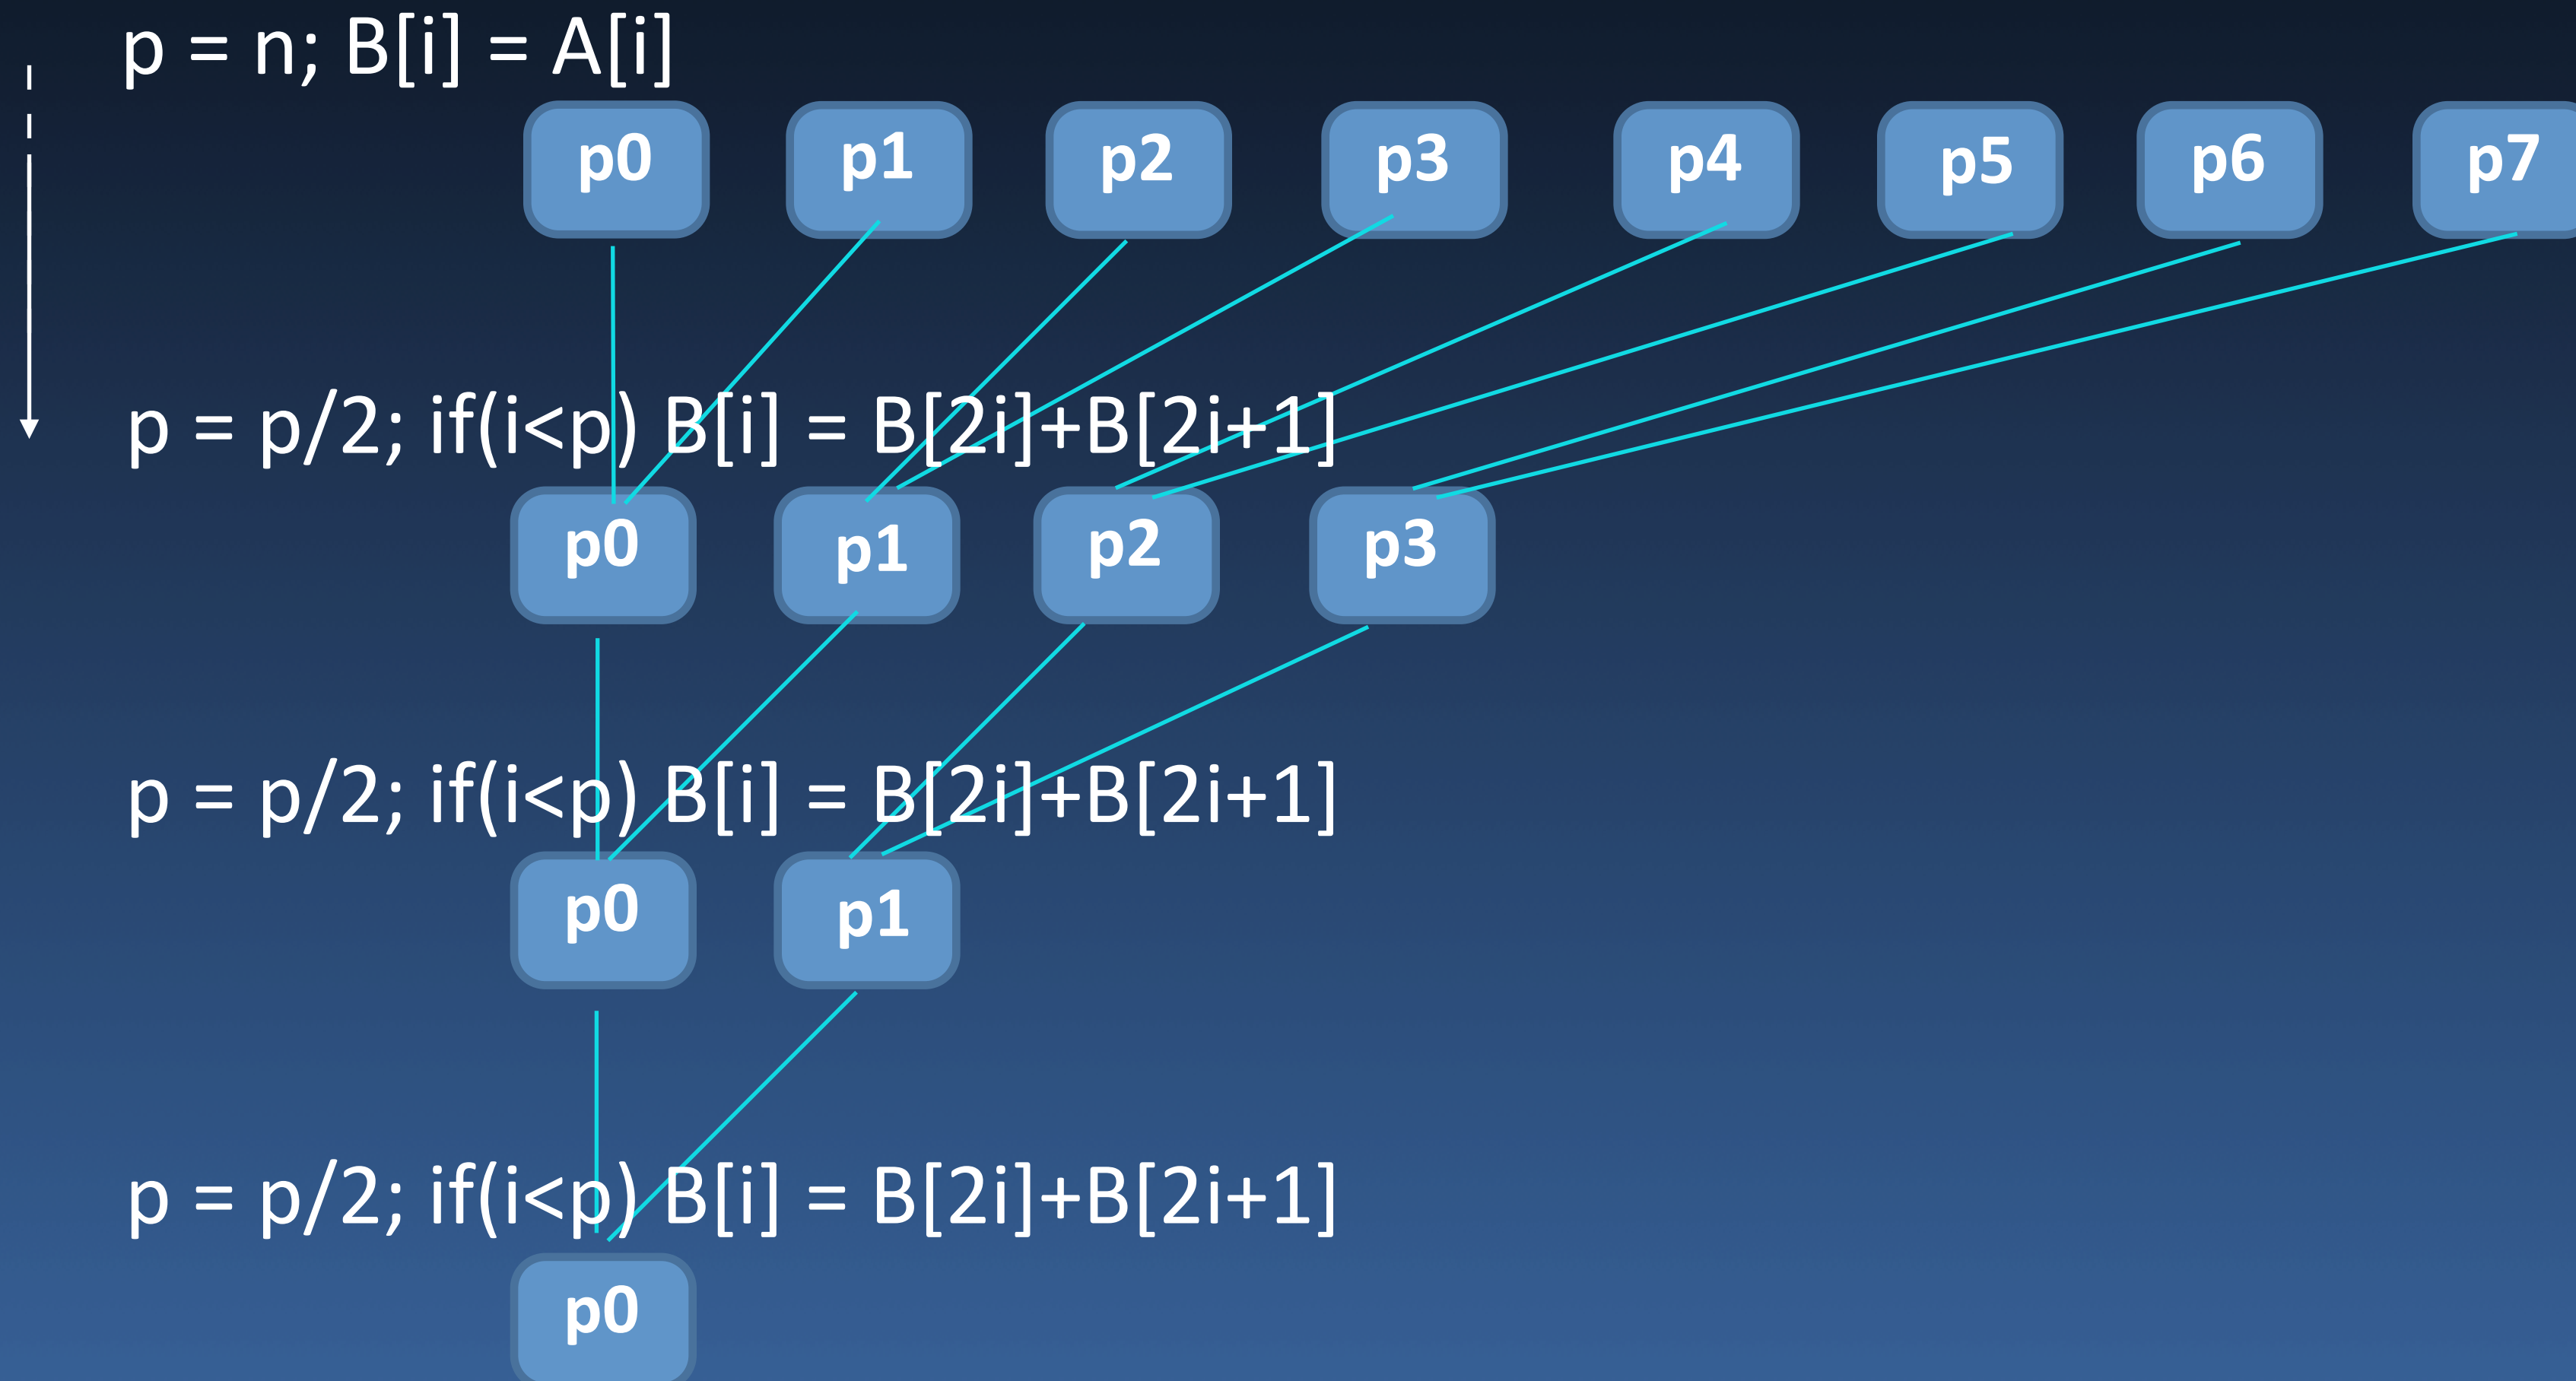
- Priority CW
 - ➔ Higher priority processor (normally lower index) wins
- Common CW
 - ➔ Succeeds only if all writes have the same value
- Arbitrary/Random CW
 - ➔ One of the values is randomly chosen

EREW ≤ CREW ≤ Common ≤ Arbitrary ≤ Priority


Less powerful

More powerful

Parallel Addition



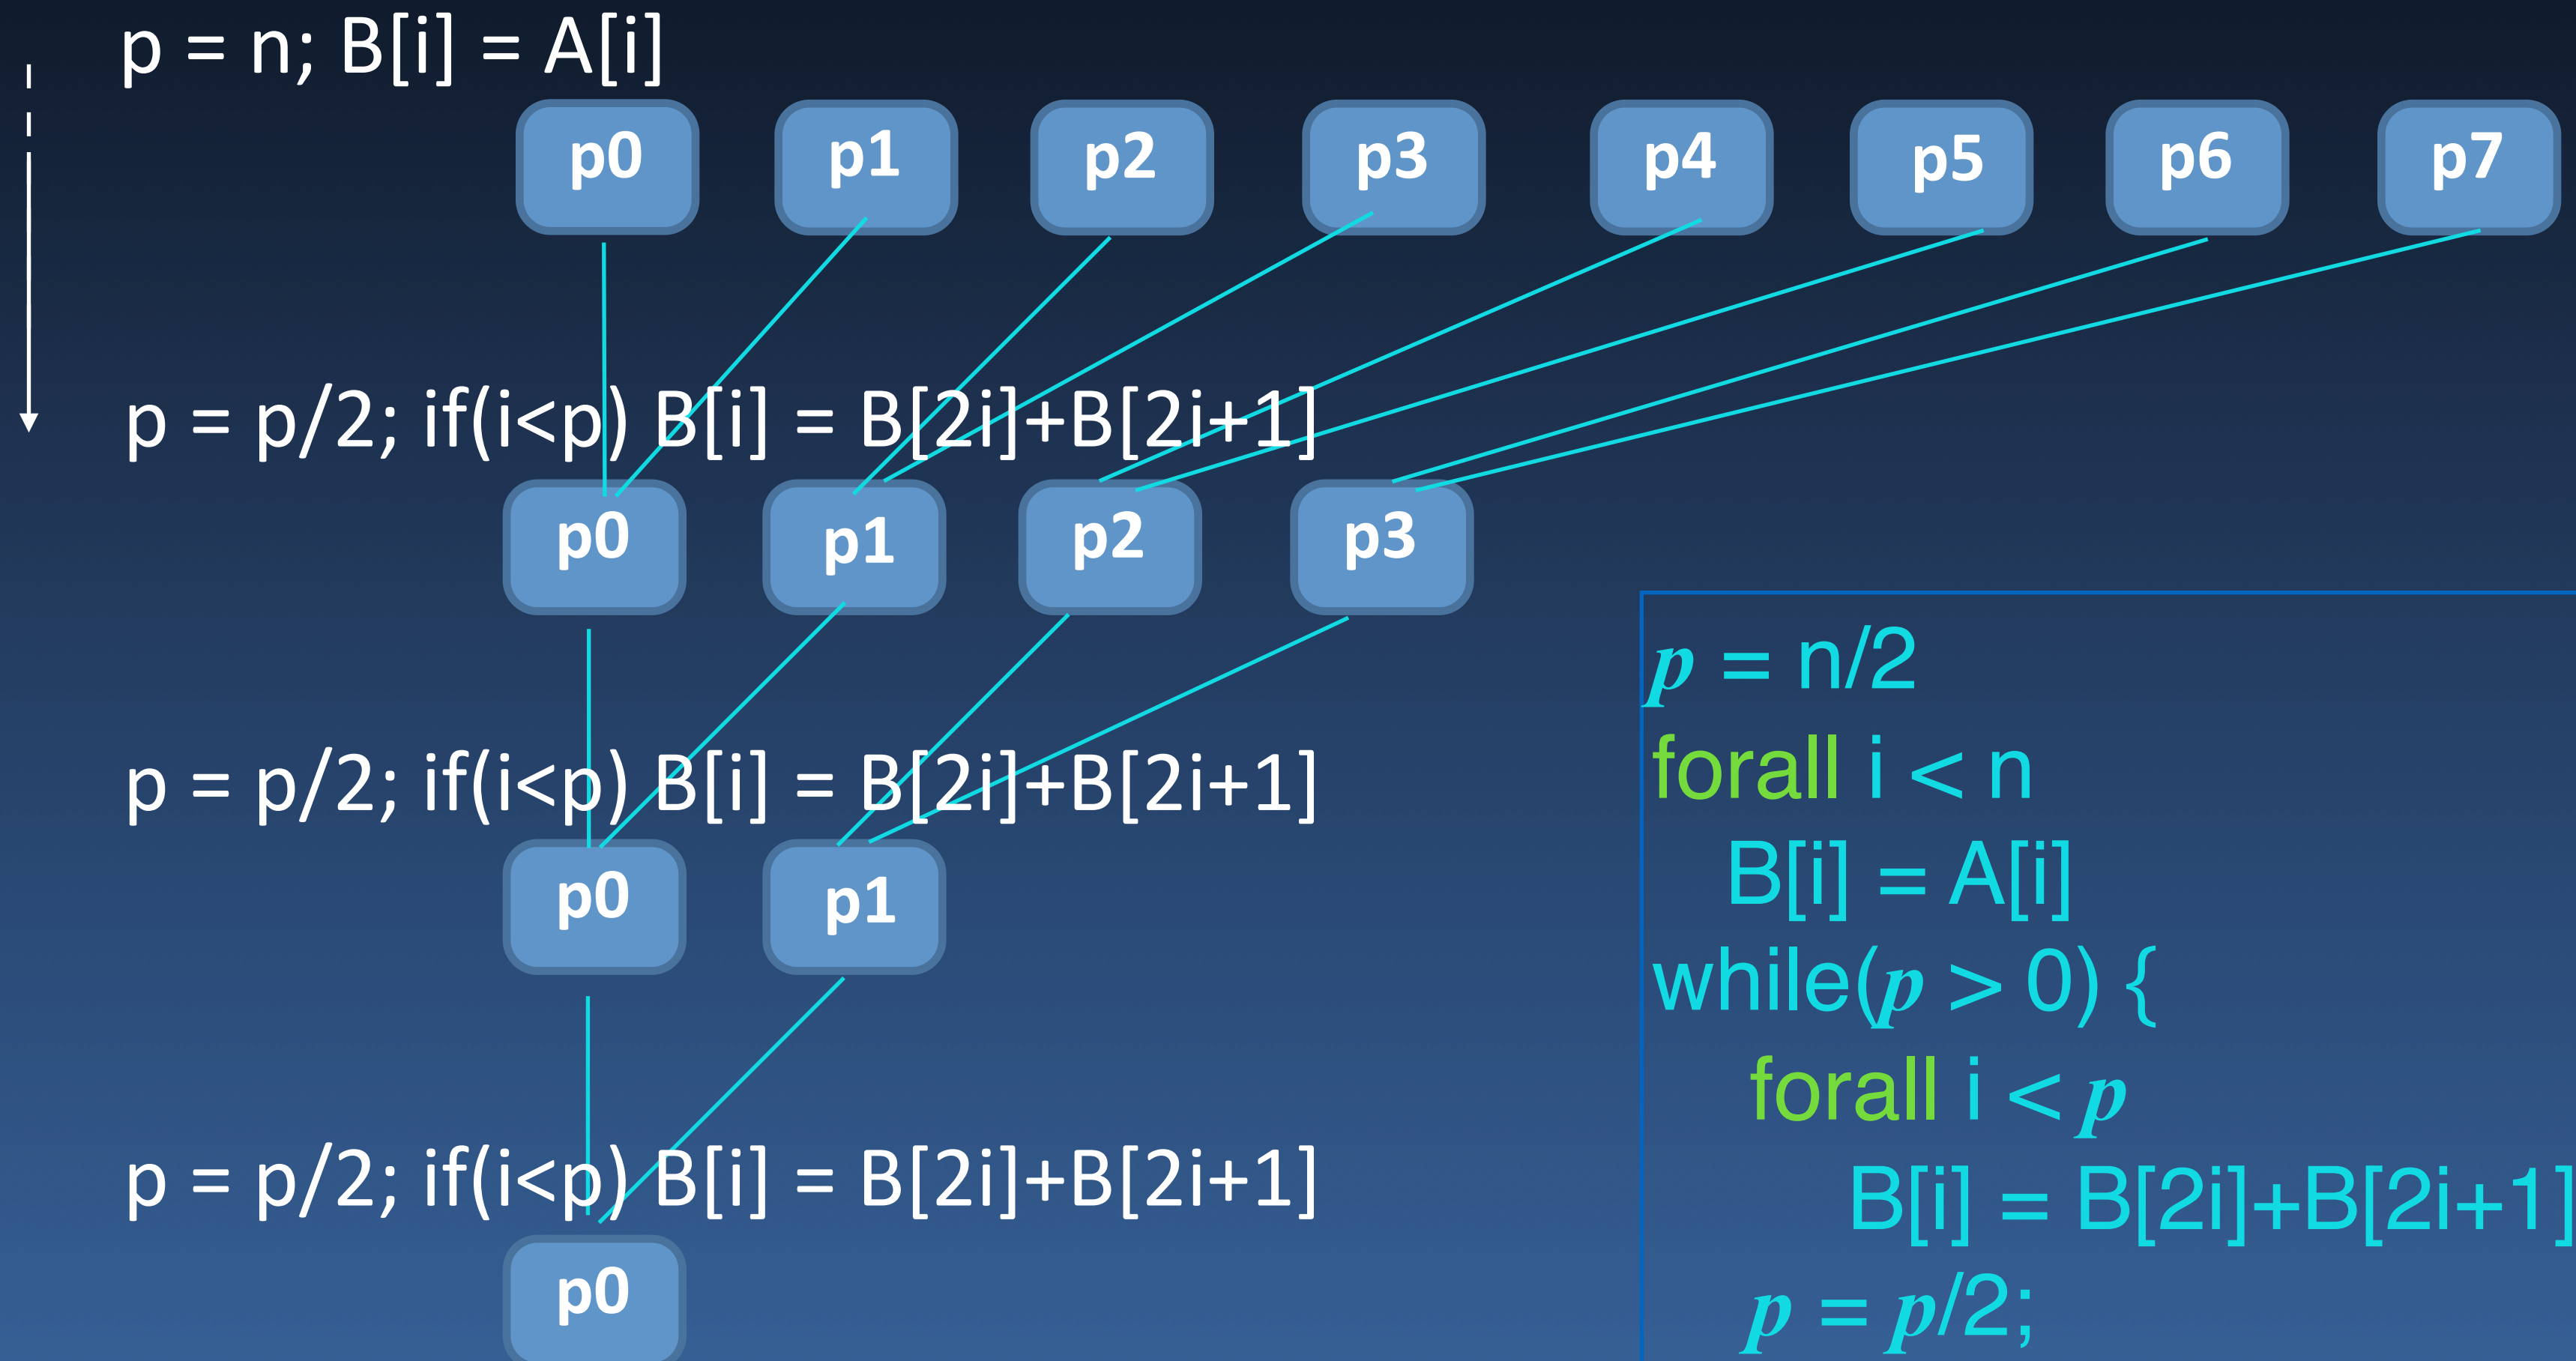
Parallel Addition



```
p = n/2
forall i < n
    B[i] = A[i]
while(p > 0) {
    forall i < p
        B[i] = B[2i] + B[2i+1]
    p = p/2;
}
```

(assumes *n* is a power of 2)

Parallel Addition



```
p = n/2  
forall i < n  
    B[i] = A[i]  
while(p > 0) {  
    forall i < p  
        B[i] = B[2i] + B[2i+1]  
    p = p/2;  
}
```

(assumes *n* is a power of 2)

- processors: *n*
- time: $O(\log n)$
- Speed-up: $n/(\log n)$
- Efficiency: $1/\log(n)$
- Cost: $n \log n$
- Work: *n*

Linear Search

$p < n$

- n input integers in n memory cells
- Does x exist in the input?
 - x is initially stored in a known shared memory location

Algorithm

step1: If p_0 , answer = 0; broadcast n, x

step2: $\forall p_i$: search in i^{th} $[n/p\text{-size}]$ block and {set flag f_i }

step3: If p_0 , check if {any} flag is 1, and print answer

EREW

• $\log(p)$

• n/p

• $\log(p)$

CREW

• 1

• n/p

• $\log(p)$

CRCW

• 1

• n/p

• 1

- Two parameters
 - ➔ $p(n)$, $t(n)$
- Generally, use work, $W(n)$
- If $W(n)$ similar, use $t(n)$
- Speedup/Scalability
 - ➔ Absolute: over best sequential algorithm
 - ➔ Relative: over the 1-processor implementation of the same algorithm

Performance Evaluation

- Two parameters

- $p(n)$, $t(n)$

- Generally, use work, $W(n)$

- If $W(n)$ similar, use $t(n)$

- Speedup/Scalability

- Absolute: over best sequential algorithm

- Relative: over the 1-processor implementation of the same algorithm

- Work-optimal \Rightarrow work = $O(\text{serial complexity})$

- $p(n)$ is hidden but important

- ▶ $W_1(n) = O(n)$; $t_1(n) = O(n)$

- ▶ $W_2(n) = O(n \log n)$ and $t_2(n) = O(\log n)$

Work Time Scheduling Principle

- Design algorithm in terms of
 - ➔ Total work done per 'time step': $W_i(n)$
 - ➔ $t(n)$ steps
- Total work done $W(n) = \sum W_i(n)$
- For each time step i :
 - ➔ divide the work $W_i(n)$ among p processors
 - ▶ Time $\leq \sum \lceil W_i(n)/p \rceil \leq \lfloor W(n)/p \rfloor + t(n)$
- Cost = $t(n,p) * p$

Work \leq Cost. Cost optimality is more stringent.

Brent's Theorem

- Time taken by p processors:
 - $t(n, p) = O(W(n)/p + t(n))$
- Cost = $p * t(n, p) = O(W(n) + p * t(n))$
- Work = Cost if:
 - $W(n) + p * t(n) = O(W(n))$
 - Or, $p = O(W(n)/t(n))$

- If **sequentially** optimal algorithm is $O(t'(n))$

→ Work done by **Work-optimal** parallel algorithm:

▶ $O(t'(n))$ (with time $t(n)$).

→ Work-scheduling on p processors takes time:

▶ $t(n, p) = O(t'(n)/p + t(n))$

→ Optimal speed-up: $t'(n)/t(n, p) = \theta(p)$, if

▶ $[p \cdot t'(n)] / [t'(n) + p \cdot t(n)] = \theta(p)$

- **Work-time** optimal if:

→ $t(n)$ cannot be improved

Why PRAM?

- Easy to design, specify, analyze algorithms
 - ➔ Independent of machine details
- Fidelity of predicted performance
 - ➔ Not many surprises for shared-memory architecture
 - ➔ Partly successful for distributed memory
 - ▶ Memory-access and message latency can often be bounded
- Strong model
 - ➔ Possible to simulate on a wide variety of hardware
 - ➔ Poor PRAM solution often implies a hard problem

Why PRAM?

- Easy to design, specify, analyze algorithms
 - ➔ Independent of machine details
- Fidelity of predicted performance
 - ➔ Not many surprises for shared-memory architecture
 - ➔ Partly successful for distributed memory
 - ▶ Memory-access and message latency can often be bounded
- Strong model
 - ➔ Possible to simulate on a wide variety of hardware
 - ➔ Poor PRAM solution often implies a hard problem

- Fine-grained synchronization

- PRAM model
 - EREW, CREW, CRCW variants
- Work-Time scheduling principle
- Work and Work-time optimality