

COL380

Introduction to  
Parallel & Distributed Programming

# Agenda

- MPI

# Starting and Ending

- `MPI_Init(&argc, &argv);`

`MPI_Init_thread`

→ Needed before any other MPI call

```
int nump, id;  
MPI_Comm_size (MPI_COMM_WORLD, &nump);  
MPI_Comm_rank (MPI_COMM_WORLD, &id);
```

- `MPI_Finalize();`

→ Required

```
int MPI_Send(void* buf, int count, MPI_Datatype datatype, int dest,  
            int tag, MPI_Comm comm)
```

- message contents    block of memory
- count                number of items in message
- message type        MPI\_Datatype of each item
- destination         rank of recipient
- tag                  integer “message identifier”
- communicator



## Send/Receive

```
int MPI_Send(void* buf, int count, MPI_Datatype datatype, int dest,  
            int tag, MPI_Comm comm)
```

```
int MPI_Recv(void* buf, int count, MPI_Datatype datatype, int  
            source, int tag, MPI_Comm comm, MPI_Status *status)
```

• block of memory	• message contents	memory buffer to store received message
• number of items in message	• count	space in buffer, overflow error if too small
• MPI_Datatype of each item	• message type	type of each item
• rank of recipient	• source	sender's rank (or <b>MPI_ANY_SOURCE</b> )
• integer "message identifier"	• tag	message identifier (or <b>MPI_ANY_TAG</b> )
	• communicator	
	• status	information about message received

## Blocking calls

## Send/Receive

```
int MPI_Send(void* buf, int count, MPI_Datatype datatype, int dest,  
            int tag, MPI_Comm comm)
```

```
int MPI_Recv(void* buf, int count, MPI_Datatype datatype, int  
            source, int tag, MPI_Comm comm, MPI_Status *status)
```

• block of memory	• message contents	memory buffer to store received message
• number of items in message	• count	space in buffer, overflow error if too small
• MPI_Datatype of each item	• message type	type of each item
• rank of recipient	• source	sender's rank (or <b>MPI_ANY_SOURCE</b> )
• integer "message identifier"	• tag	message identifier (or <b>MPI_ANY_TAG</b> )
	• communicator	
	• status	information about message received

## Blocking calls

## Send/Receive

```
int MPI_Send(void* buf, int count, MPI_Datatype datatype, int dest,  
            int tag, MPI_Comm comm)
```

MATCHING (Per context)

```
int MPI_Recv(void* buf, int count, MPI_Datatype datatype, int  
            source, int tag, MPI_Comm comm, MPI_Status *status)
```

• block of memory

• number of items in message

• MPI\_Datatype of each item

• rank of recipient

• integer “message identifier”

• message contents

• count

• message type

• source

• tag

• communicator

• status

• memory buffer to store received message

• space in buffer, overflow error if too small

• type of each item

• sender’s rank (or MPI\_ANY\_SOURCE)

• message identifier (or MPI\_ANY\_TAG)

• information about message received



# Eager vs Rendezvous

## Eager

- Send-stub packetizes and transmits  
(May save a local message copy)
- Send-stub signals Done
- Recv-stub continuously accepts
- Delivered when Recv call matches

## Rendezvous

- Send-stub transmits envelope info  
(May save local message copy)
- Recv-stub continuously accepts envelope info
- Recv-stub may signal OK (if it has space)  
Or, wait for matching Recv call to be made
- Recv-stub sets up “RDMA” with Send-stub
- Data transmitted
- Recv-stub signals Done
- Send-stub signals Done



# Send/Recv Synchronization

- **Blocking**

- ➔ Send returns after some progress guarantee
  - ▶ Receive completed?
  - ▶ Synchronization (up to network delay)

- **Immediate**

- ➔ Send returns with no progress guarantee
- ➔ Receiver may also proceed immediately (message arrives later)

- **Standard mode:**
  - ➔ implementation dependent
- **Buffered mode**
  - ➔ MPI saves a copy of message, Receiver can post later
  - ➔ User provided buffer
- **Synchronous mode**
  - ➔ Will complete only once a matching receive has started
- **Ready mode**
  - ➔ Send may start only if a matching receive has already been called
  - ➔ Helps performance

## Send Semantics

- **Standard mode:**
    - ➔ implementation dependent
  - **Buffered mode**
    - ➔ MPI saves a copy of message, Receiver can post later
    - ➔ User provided buffer
  - **Synchronous mode**
    - ➔ Will complete only once a matching receive has started
  - **Ready mode**
    - ➔ Send may start only if a matching receive has already been called
    - ➔ Helps performance
- **MPI\_Send/MPI\_Recv are blocking**
    - ➔ Recv blocks until output buffer is filled
    - ➔ Send blocks until some 'progress'



## Send Semantics

- **Standard mode:** `MPI_Send`
    - ➔ implementation dependent
  - **Buffered mode** `MPI_Bsend`
    - ➔ MPI saves a copy of message, Receiver can post later
    - ➔ User provided buffer `See MPI_Buffer_attach`
  - **Synchronous mode** `MPI_Ssend`
    - ➔ Will complete only once a matching receive has started
  - **Ready mode** `MPI_Rsend`
    - ➔ Send may start only if a matching receive has already been called
    - ➔ Helps performance
- `MPI_Send/MPI_Recv` are blocking
    - ➔ Recv blocks until output buffer is filled
    - ➔ Send blocks until some 'progress'



- MPI Blocking Send- Receive
  - ➔ Semantics
  - ➔ Matching