

COL380

Introduction to
Parallel & Distributed Programming

Agenda

- MPI non-blocking
- MPI datatypes
- MPI collectives

- In order (per pair and tag)
 - ➔ Multi-threaded applications need to be coordinate
- Progress
 - ➔ For a matching send/Recv pair, at least one of these two will complete
- Fairness not guaranteed
 - ➔ A Send or a Recv may starve because all matches are satisfied by others
- Resource limitation can cause deadlocks
- Ready/Synchronous sends requires the least resources
 - ➔ Also used for debugging

Example

If (rank == 0):

Send(sbuffer0, to 1);

Recv(rbuffer0, from 1);

else:

Send(sbuffer1, to 0);

Recv(rbuffer1, from 0);

Deadlock

if neither send can copy out its sbuffer

Non-blocking Call

```
int MPI_Isend(void* buf, int count, MPI_Datatype datatype, int dest, int tag,  
MPI_Comm comm, MPI_Request *request)
```

```
int MPI_Irecv(void* buf, int count, MPI_Datatype datatype, int source, int  
tag, MPI_Comm comm, MPI_Request *request)
```

Non-blocking calls

Returns even before buf copied out; caller must not use.

Example

If (rank == 0):

MPI_Isend(sbuffer0, to 1);

MPI_Irecv(rbuffer0, from 1);

Wait for earlier calls to finish

Will not deadlock

else:

MPI_Isend(sbuffer1, to 0);

MPI_Irecv(rbuffer1, from 0);

Wait for earlier calls to finish

- **MPI_Wait**(&request, &status)

- ➔ status similar to MPI_recv
- ➔ Blocks as per the blocking version's semantics
 - ▶ Send: message was copied out, Recv was started, etc.
 - ▶ Recv: Wait for data to fill
- ➔ Request is freed as a side-effect

- **MPI_Test**(&request, &flag, &status)

- ➔ Non-blocking poll
- ➔ flag indicates whether operation is complete
- ➔ Request is freed as a side-effect

Use MPI_Request_get_status to retain request
Later MPI_Request_free

- **MPI_Wait**(&request, &status)

- ➔ status similar to MPI_recv
- ➔ Blocks as per the blocking version's semantics
 - ▶ Send: message was copied out, Recv was started, etc.
 - ▶ Recv: Wait for data to fill
- ➔ Request is freed as a side-effect

Also see:

MPI_Waitany, MPI_Waitall, MPI_Waitsome
MPI_Testany, MPI_Testall, MPI_Testsome

- **MPI_Test**(&request, &flag, &status)

- ➔ Non-blocking poll
- ➔ flag indicates whether operation is complete
- ➔ Request is freed as a side-effect

- **Send - Recv is point-to-point**
 - ➔ Call-to-call matching
 - ➔ Integer tag to control matching
 - ➔ Wildcard matching: MPI_ANY_SOURCE and MPI_ANY_TAG
- **Recv buffer must contain enough space for message**
 - ➔ Receiving fails otherwise
 - ➔ Can query the actual count received (MPI_Get_count)
 - ▶ Send determines the actual number sent
 - ➔ type parameters determines data structure ↔ message buffer copying

MPI Data types

• MPI_CHAR	signed char
• MPI_SHORT	signed short int
• MPI_INT	signed int
• MPI_LONG	signed long int
• MPI_LONG_LONG_INT	signed long long int
• MPI_LONG_LONG	signed long long int
• MPI_SIGNED_CHAR	signed char
• MPI_UNSIGNED_CHAR	unsigned char
• MPI_UNSIGNED_SHORT	unsigned short int
• MPI_UNSIGNED	unsigned int
• MPI_UNSIGNED_LONG	unsigned long int
• MPI_UNSIGNED_LONG_LONG	unsigned long long int
• MPI_FLOAT	float
• MPI_DOUBLE	double
• MPI_LONG_DOUBLE	long double
• MPI_WCHAR	wchar_t
• MPI_BYTE	

Objects of type
MPI_Datatype

- MPI does not understand language's layout (struct, e.g.)

- ➔ Too system architecture dependent

MPI_INT, MPI_FLOAT ..

- Typemap:

- ➔ (type_0, disp_0), ..., (type_n, disp_n)

- ➔ i^{th} entry is of type i and starts at byte base + disp_ i

- MPI does not understand language's layout (struct, e.g.)

- ➔ Too system architecture dependent

MPI_INT, MPI_FLOAT ..

- Typemap:

- ➔ (type_0, disp_0), ..., (type_n, disp_n)

- ➔ i^{th} entry is of type i and starts at byte base + disp_ i

```
MPI_Datatype newtype;
```


```
MPI_Type_contiguous(count, MPI_INT, &newtype);
```


Blocks

- Equally-spaced blocks of the known datatype

→ `MPI_Type_vector`(³blockcount, ²blocklength, ⁴blockstride, knowntype, &newtype);


- ▶ Assume contiguous copies of 'knowntype'
- ▶ Stride between blocks specified in units of knowntype
- ▶ All picked blocks are of the same length



→ `MPI_Type_create_hvector`(blk_count, blk_length, bytestride, knowntype, &newtype);
Gap between blocks is in bytes

- `MPI_Type_indexed`(⁵count, ^{2,1,1,1,2}array_of_blocklengths,
^{0,4,6,8,10}array_of_offsets, knowntype, &newtype);

- Blocks can contain different number of copies
- And may have different strides
- But the same data type



Struct

- **MPI_Type_create_struct**(count, array_of_blocklengths, array_of_byteoffsets, array_of_knowntypes, &newtype)

→ Example:

- ▶ Suppose Type0 = {(double, 0), (char, 8)},
- ▶ int BL[] = {2, 1, 3}, Disp[] = {0, 16, 26};
- ▶ MPI_Datatype Typ[] = {MPI_FLOAT, Type0, MPI_CHAR}

→ MPI_Type_create_struct(3, BL, Disp, Typ, &newtype):


- ▶ (float, 0), (float, 4), (double, 16), (char, 24), (char, 26), (char, 27), (char, 28)



- **MPI_Type_create_struct**(count, array_of_blocklengths, array_of_byteoffsets, array_of_knowntypes, &newtype)

→ Example:

▶ Suppose Type0 = {(double, 0), (char, 8)},



▶ int BL[] = {2, 1, 3}, Disp[] = {0, 16, 26};



▶ MPI_Datatype Typ[] = {MPI_FLOAT, Type0, MPI_CHAR}


→ MPI_Type_create_struct(3, BL, Disp, Typ, &newtype):

▶ (float, 0), (float, 4), (double, 16), (char, 24), (char, 26), (char, 27), (char, 28)

- **MPI_Type_create_struct**(count, array_of_blocklengths, array_of_byteoffsets, array_of_knowntypes, &newtype)

→ Example:

▶ Suppose Type0 = {(double, 0), (char, 8)},



▶ int BL[] = {2, 1, 3}, Disp[] = {0, 16, 26};



▶ MPI_Datatype Typ[] = {MPI_FLOAT, Type0, MPI_CHAR}

→ MPI_Type_create_struct(3, BL, Disp, Typ, &newtype):

▶ (float, 0), (float, 4), (double, 16), (char, 24), (char, 26), (char, 27), (char, 28)

MPI_Type_get_contents(..)

- `MPI_Type_commit(&datatype)`
 - ➔ A datatype object must be committed before communication
- `MPI_Type_size(datatype, &size)`
 - ➔ Total size in bytes
- `MPI_Type_get_extent(datatype, &beg, &extent);`
- `MPI_Type_create_resized(datatype, beg, extent, &newtype);`
- `MPI_Get_address(data, &Address[0]);`
- `MPI_BOTTOM`

Data Type Functions

- `MPI_Type_commit(&datatype)`

➔ A datatype object must be committed before communication

- `MPI_Type_size(datatype, &size)`

➔ Total size in bytes

- `MPI_Type_get_extent(datatype, &b`

```
MPI_Datatype atype;  
MPI_Type_contiguous(4, MPI_CHAR, &atype);  
int asize;  
MPI_Type_size(atype, &asize);  
MPI_Type_commit(&atype);  
MPI_Send(buf, nItems, atype, dest, ..);  
MPI_Recv(...);
```

- `MPI_Type_create_resized(datatype, beg, extent, &newtype);`

- `MPI_Get_address(data, &Address[0]);`

- `MPI_BOTTOM`

Derived Type Example

sendParticles(struct Particle particle[], int N):

```
MPI_Datatype Particletype;
```

```
MPI_Datatype types[3] = {MPI_INT, MPI_DOUBLE, MPI_CHAR};
```

```
int blockcount[3] = {1, 6, 7};
```

```
/* compute displacements of structure components */
```

```
MPI_Aint disp[3];
```

```
MPI_Address(particle, disp);
```

```
MPI_Address(particle[0].d, disp+1);
```

```
MPI_Address(particle[0].b, disp+2);
```

```
for (int i=2; i >= 0; i--) disp[i] -= disp[0];
```


```
MPI_Type_struct(3, blockcount, disp, types, &Particletype);
```

```
MPI_Type_commit( &Particletype);
```

```
MPI_Send(particle, N, Particletype, dest, tag, comm);
```

```
struct Particle
{
    int class;    // particle class
    double d[6]; // particle coordinates
    char b[7];    // some additional info
};
```


Data Transfer



```
MPI_Status status;
MPI_Datatype column;
MPI_Type_vector(n, 1, n+1, MPI_INT, &column);
MPI_Type_commit(&column);
if(rank == 0) {
    MPI_Send(A+n-1, 1, column, 1, tag, MPI_COMM_WORLD);
    MPI_Recv(A+n, 1, column, 1, tag, MPI_COMM_WORLD, &status);
}
if(rank == 1) {
    MPI_Recv(A, 1, column, 0, tag, MPI_COMM_WORLD, &status);
    MPI_Send(A+1, 1, column, 0, tag, MPI_COMM_WORLD);
}
```

- MPI_Barrier
 - Barrier synchronization across all members of a group
- MPI_Bcast
 - Broadcast from one member to all members of a group
- MPI_Scatter, MPI_Gather, MPI_Allgather
 - Gather data from all members of a group to one
- MPI_Alltoall
 - complete exchange or all-to-all
- MPI_Reduce, MPI_Allreduce,
 - Reduction operations
- MPI_Reduce_Scatter
 - Combined reduction and scatter operation
- MPI_Scan, MPI_Exscan
 - Prefix

- Synchronization of the calling processes
 - the call blocks until all of the processes have placed the call

```
MPI_Barrier(comm) ;
```

Broadcast

MPI_Bcast(**mesg**, **count**, **MPI_INT**, **root**, **comm**) ;

pointer on all number & type identified sender can be intercommunicator

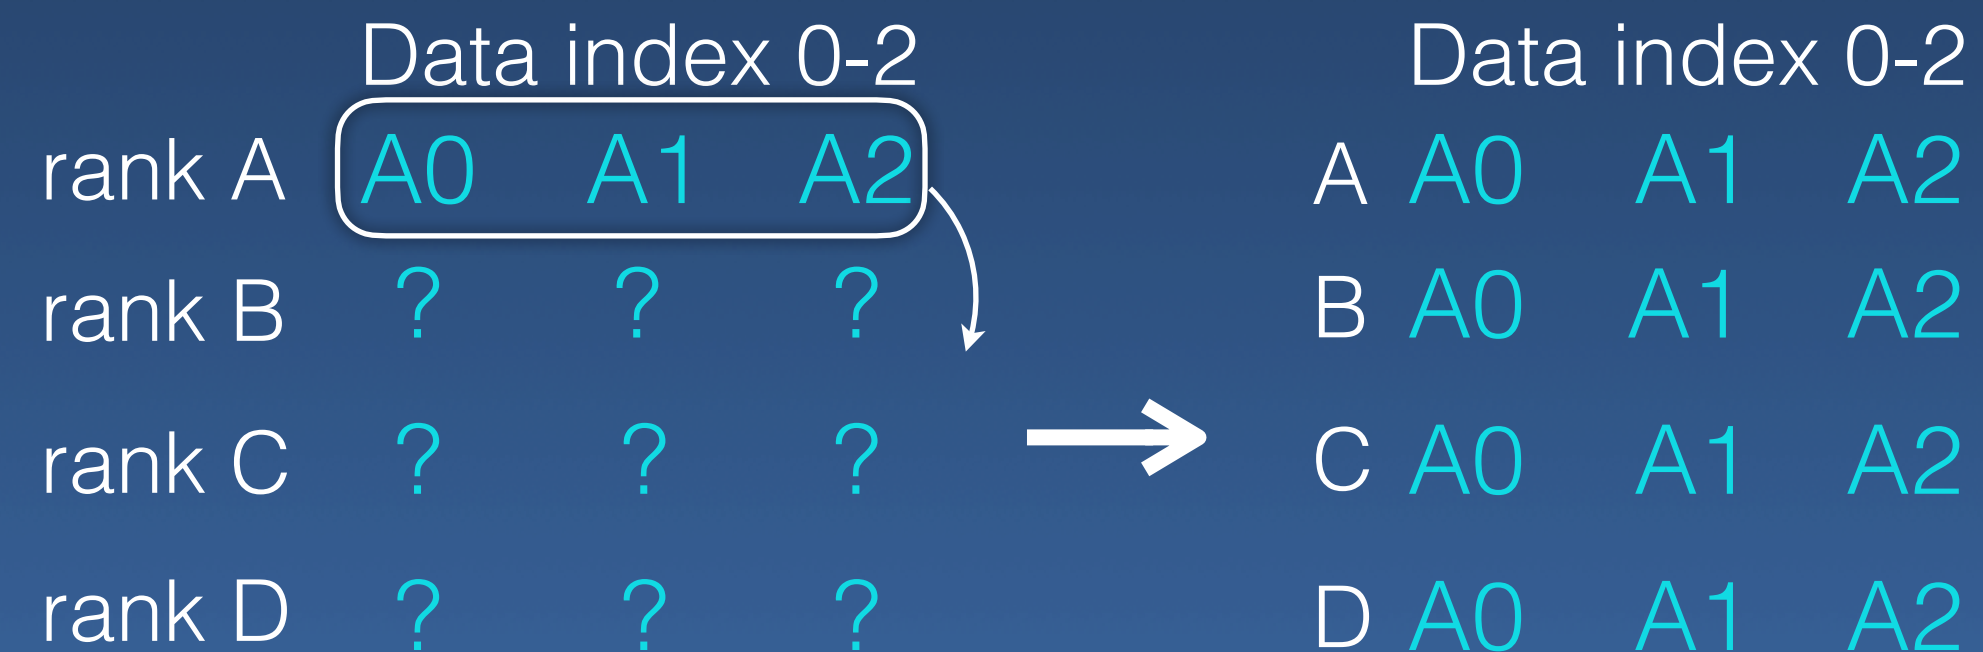
- All participants must call, match by comm and root
- No implicit synchronization

Broadcast

MPI_Bcast(**mesg**, **count**, **MPI_INT**, **root**, **comm**) ;

pointer on all number & type identified sender can be intercommunicator

- All participants must call, match by comm and root
- No implicit synchronization



- Difference in usage of blocking and non-blocking send
- Making MPI types suitable for data transfer
 - ➔ Types include 'untransferred' holes
- Introduction to group-collective calls